

Global Accelerator

Chandra Lingam

Cloud Wave LLC

Motivation

Elastic Load Balancers are regional resources

Global Accelerator – Cross-region load balancing

What is the need for cross-region load balancing?

Go Global



- Global Audience
- Performance

App

Data

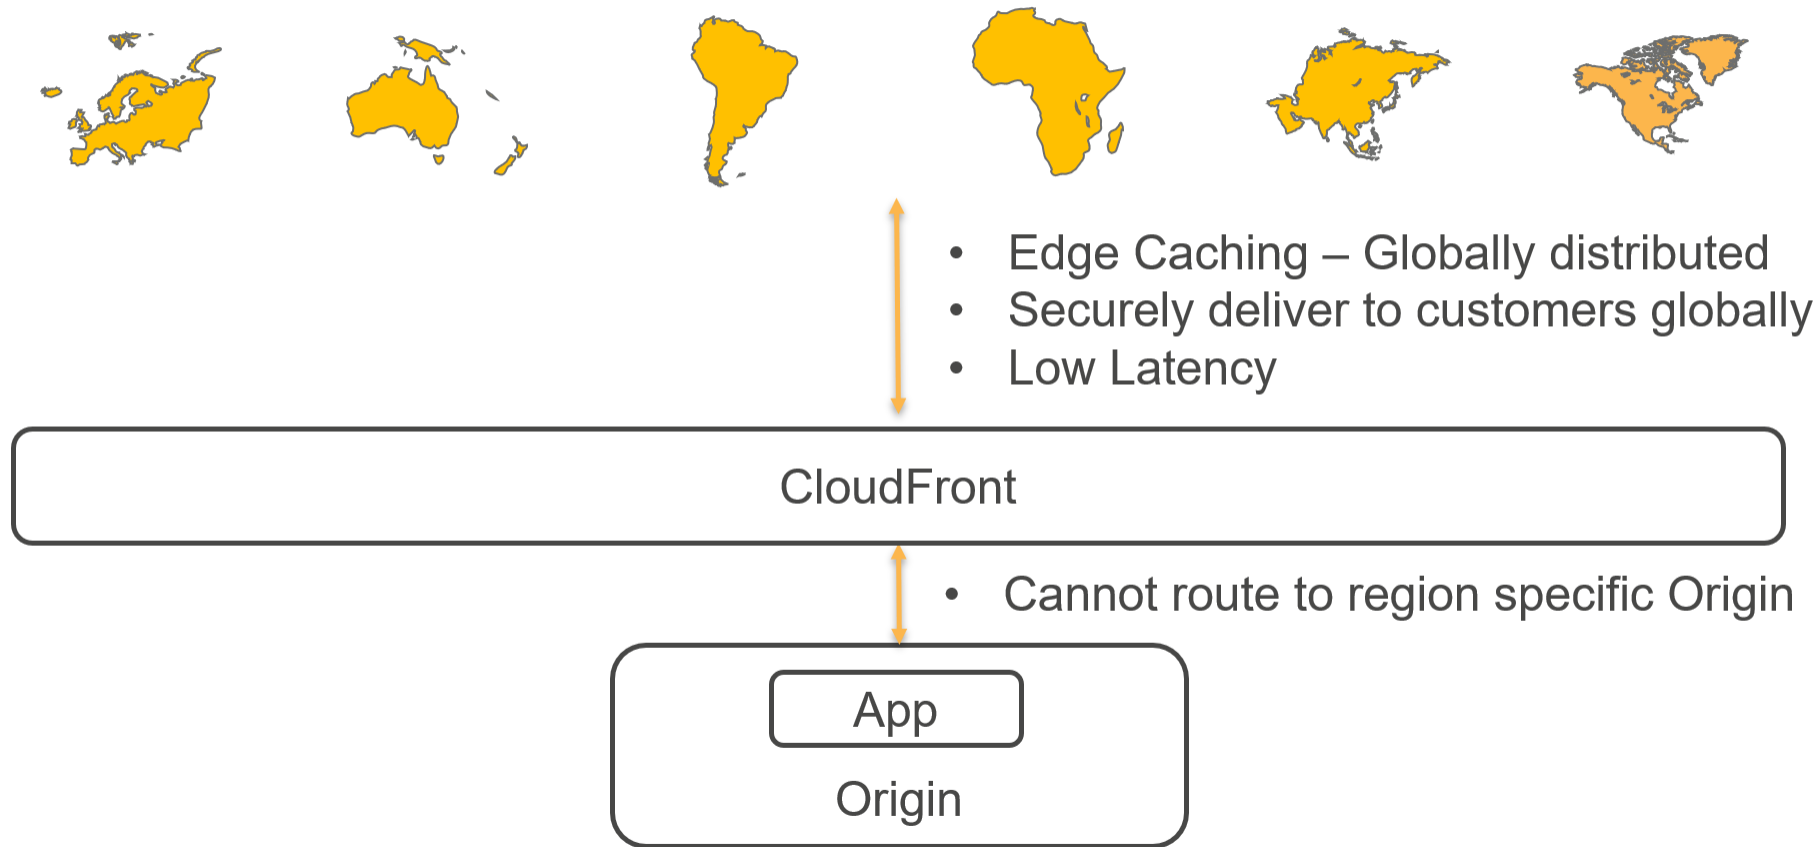
Video

Image

API

Content

CloudFront CDN



Keep your app in several regions



- Route traffic to appropriate region
- Complex setup

App
Region 1

App
Region 2

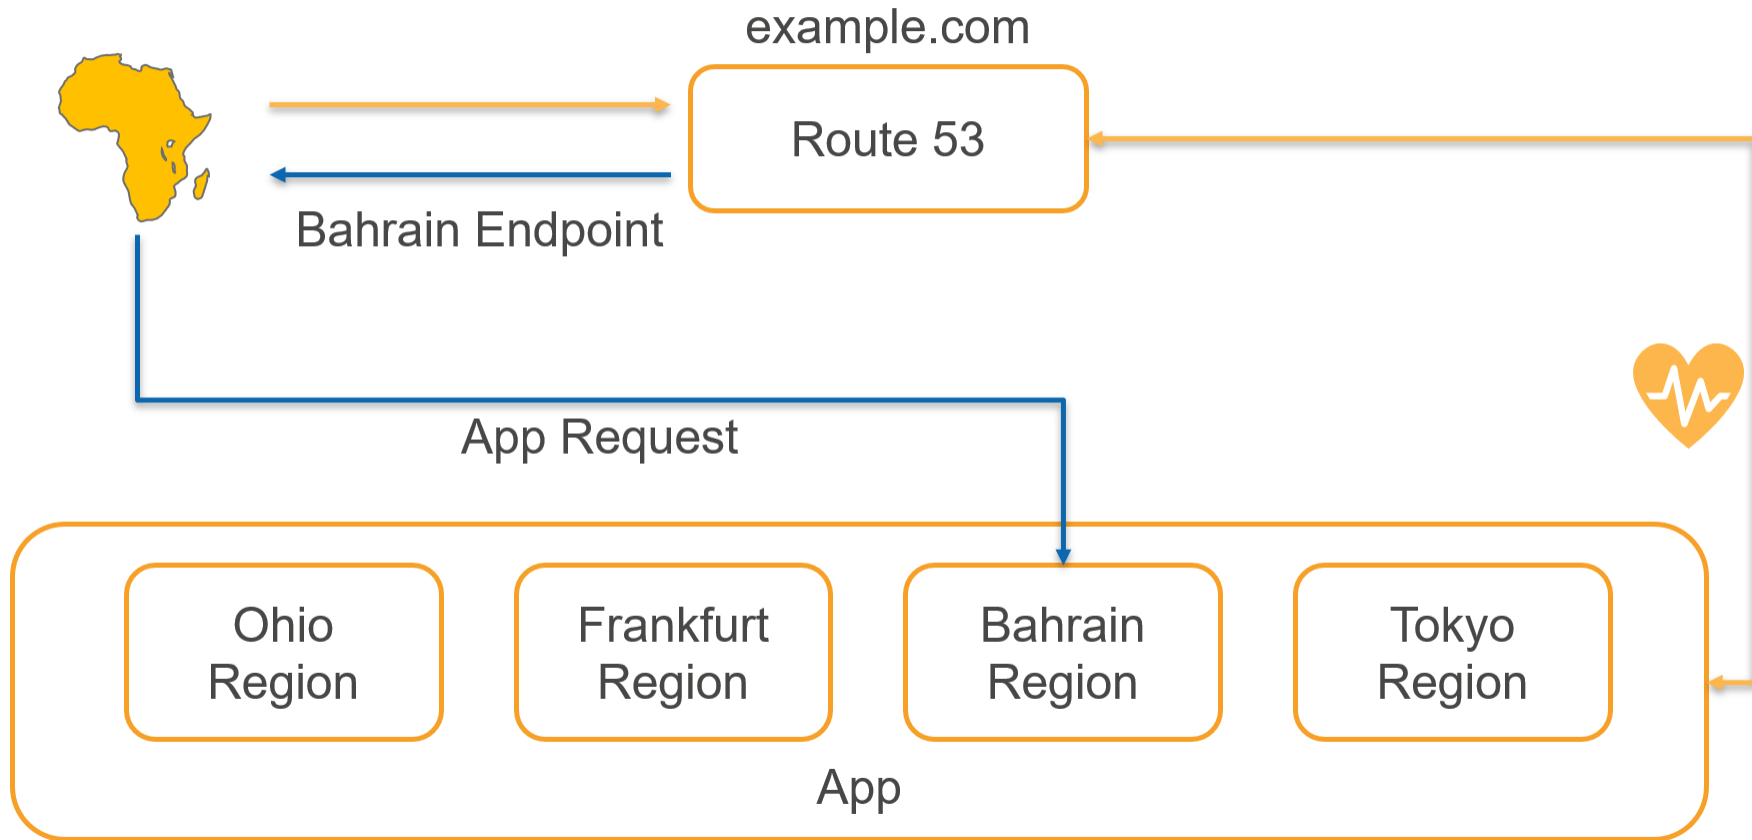
App
Region 3

App
Region 4

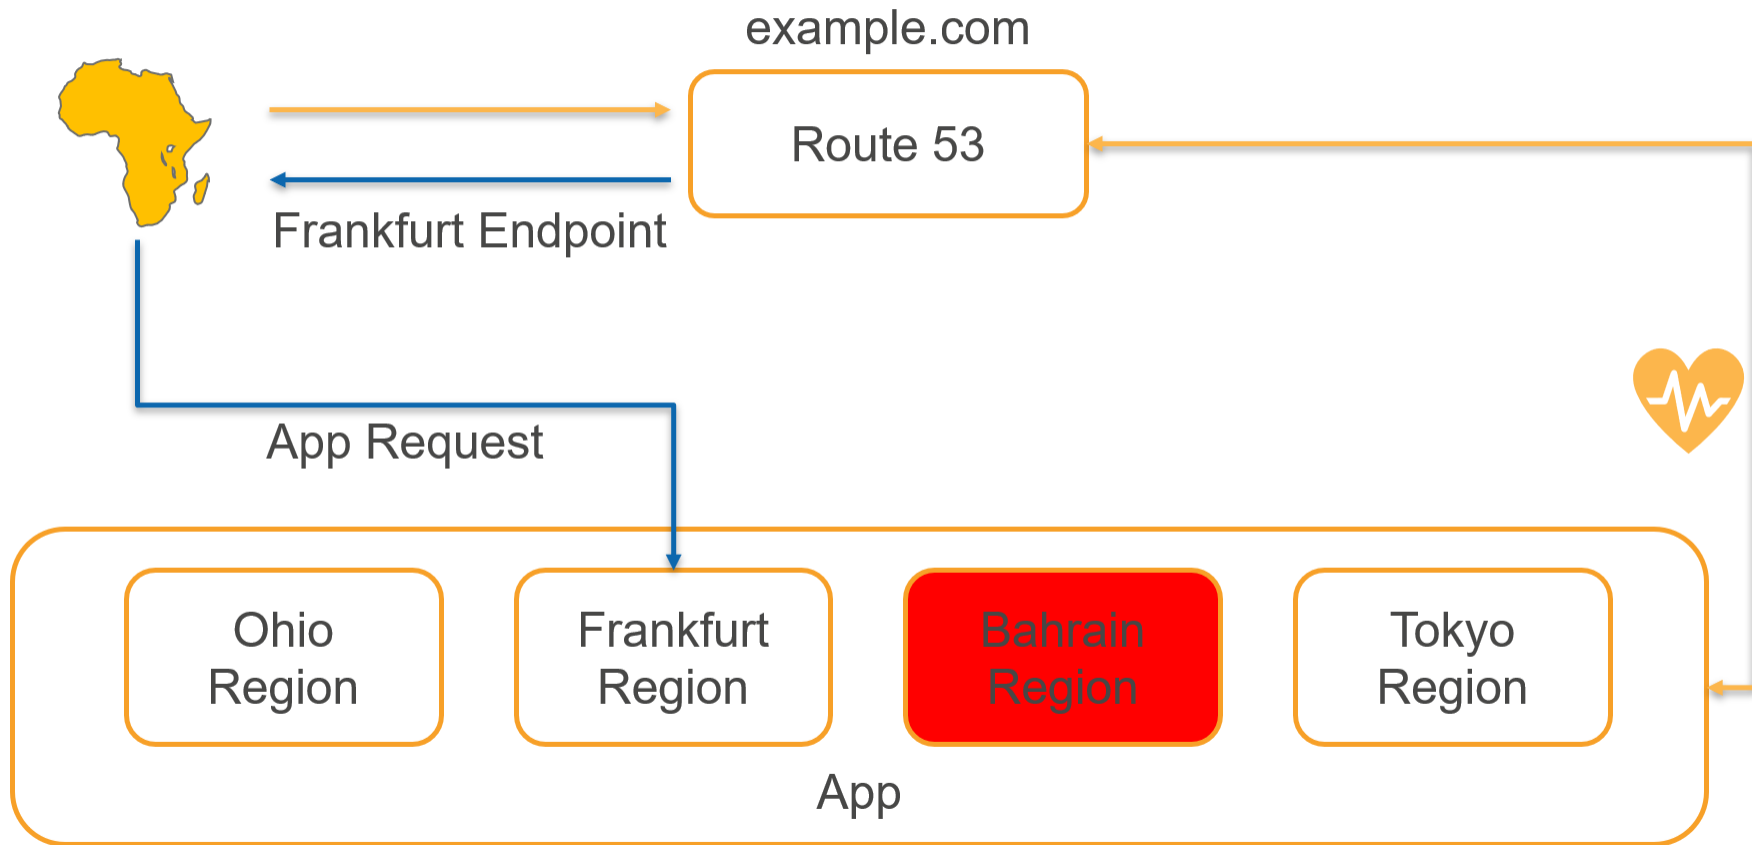
App
Region 5

App in major regions

Route 53 DNS Resolution



Route 53 DNS Resolution



Route 53 based routing issues

DNS Records are cached in client and DNS resolvers for duration specified in TTL

Delay in responding and rerouting traffic to new endpoint

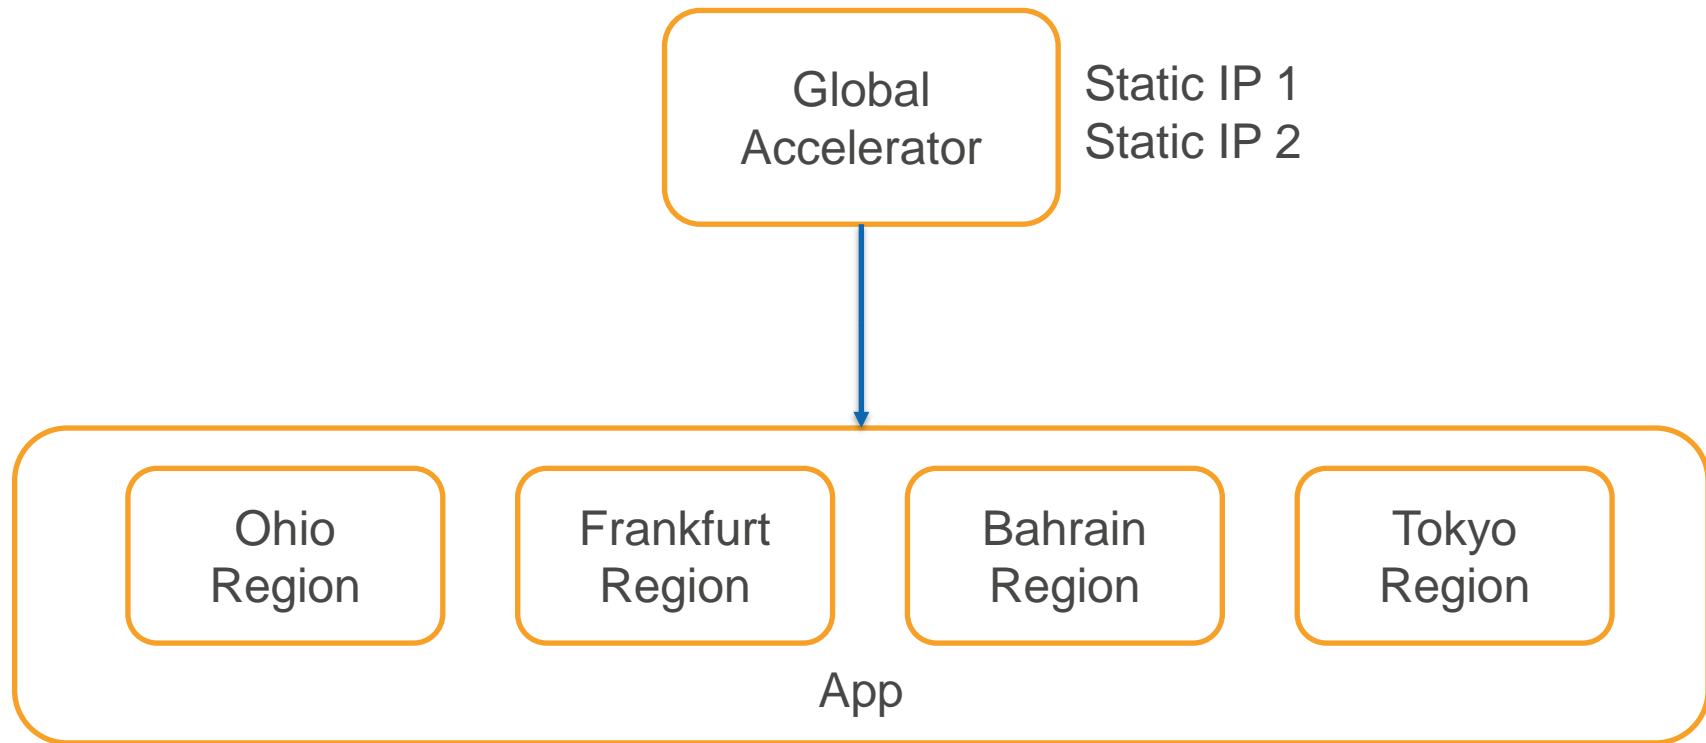
Infrastructure scaling actions results in IP address changes

- ALB adds and removes load balancer nodes

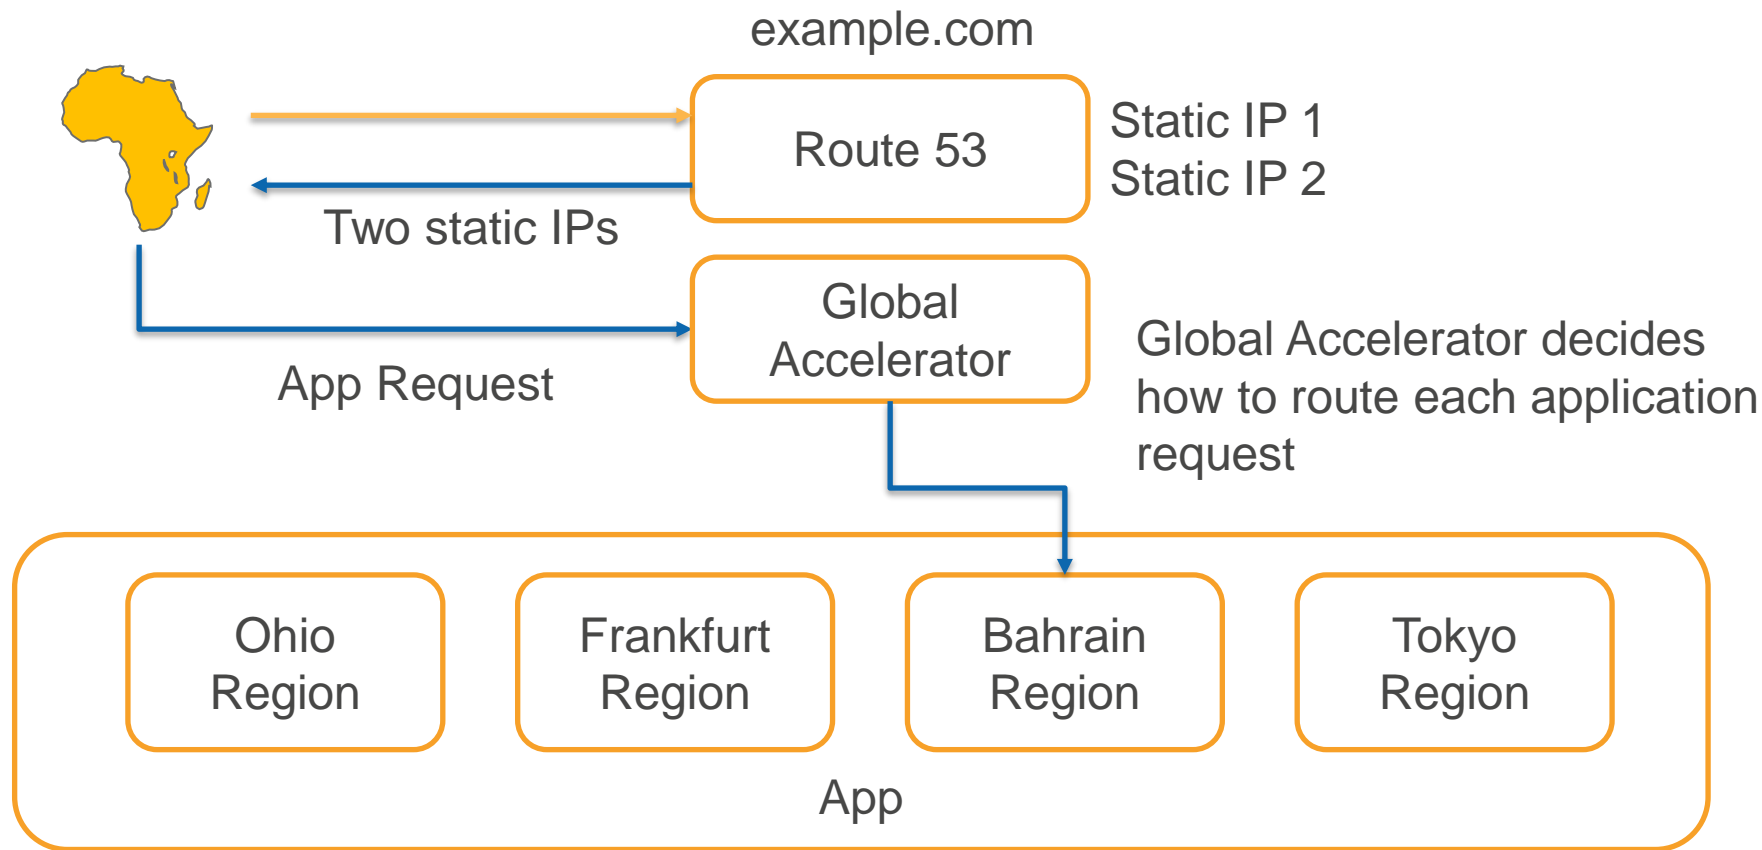
- List of load balancer nodes is maintained by ALB in Route 53

Polling based

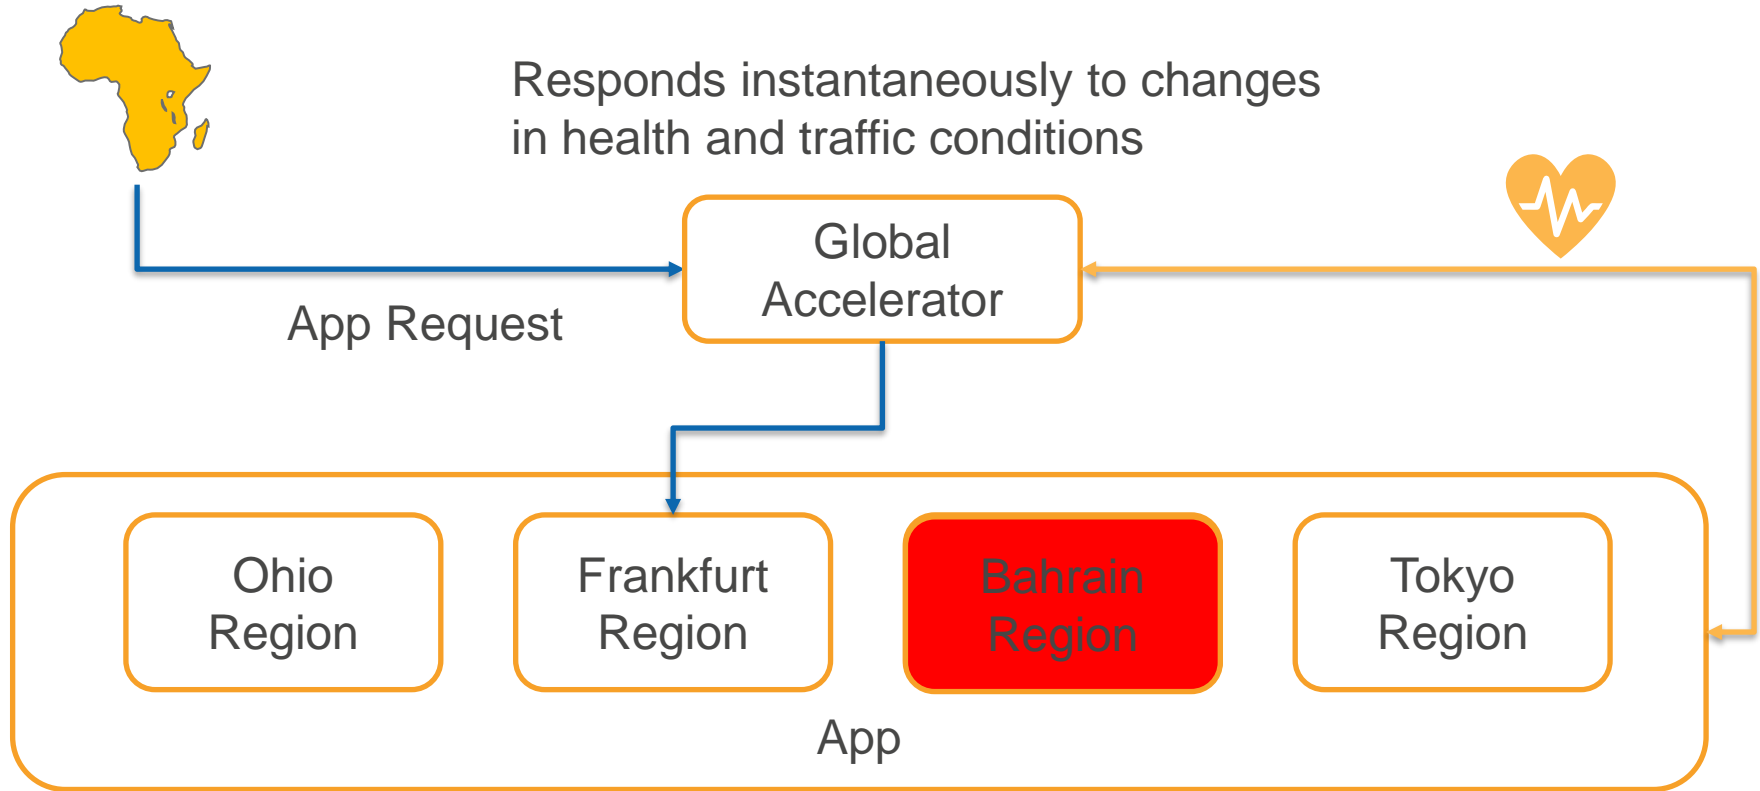
Global Accelerator



Global Accelerator



Global Accelerator



Global Accelerator Speed Test Tool

AWS Global Accelerator

Speed Comparison

About this tool

[AWS Global Accelerator](#) is a service that improves the availability and performance of your applications. This tool compares Global Accelerator to the public internet. Choose a file size to see the time to download a file from application endpoints in different AWS Regions to your browser.

Files are downloaded over HTTPS/TCP from Application Load Balancers (ALBs) in different AWS Regions to your browser. [Learn more](#)

Choose a file size and click "Start" to start the tests:

100KB ▼

Reset

We welcome suggestions for how to improve this tool. [Provide feedback](#)

ⓘ Results may differ when you run the test multiple times. Download times can vary based on factors that are external to Global Accelerator, such as the quality, capacity, and distance of the connection in the last-mile network that you're using.

Oregon (us-west-2)

Direct over internet



Total time

645ms

N. Virginia (us-east-1)

Direct over internet



Total time

146ms

AWS Global Accelerator



101ms

31% faster with AWS Global Accelerator

Ireland (eu-west-1)

Direct over internet



Total time

686ms

Global Accelerator

Clients connect to Edge location

Request travels in AWS Global network (and not public internet)

Endpoints can be private (Global Accelerator creates a peering connection)

Client-affinity – when enabled, request from a client is sent to the same endpoint

Global Accelerator

Hides infrastructure

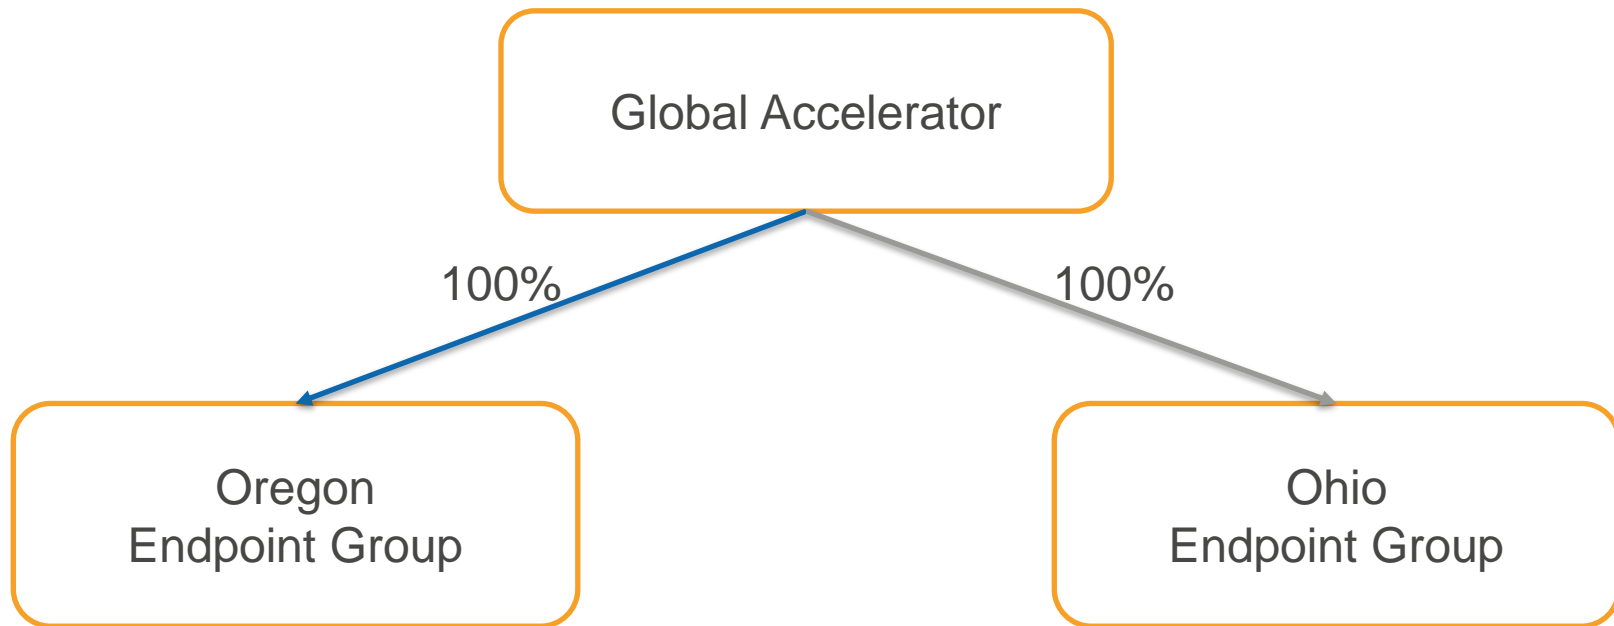
DDoS Protection

AWS Shield Standard and Advanced Protection

Bring your own IP


Traffic Dial

Global Accelerator selects a healthy endpoint group closest to the requester




Traffic Dial – Shape traffic flow between regions

50% of Oregon traffic is routed to Ohio and 50% of Ohio traffic is routed to Oregon




Traffic Dial – During Deployment

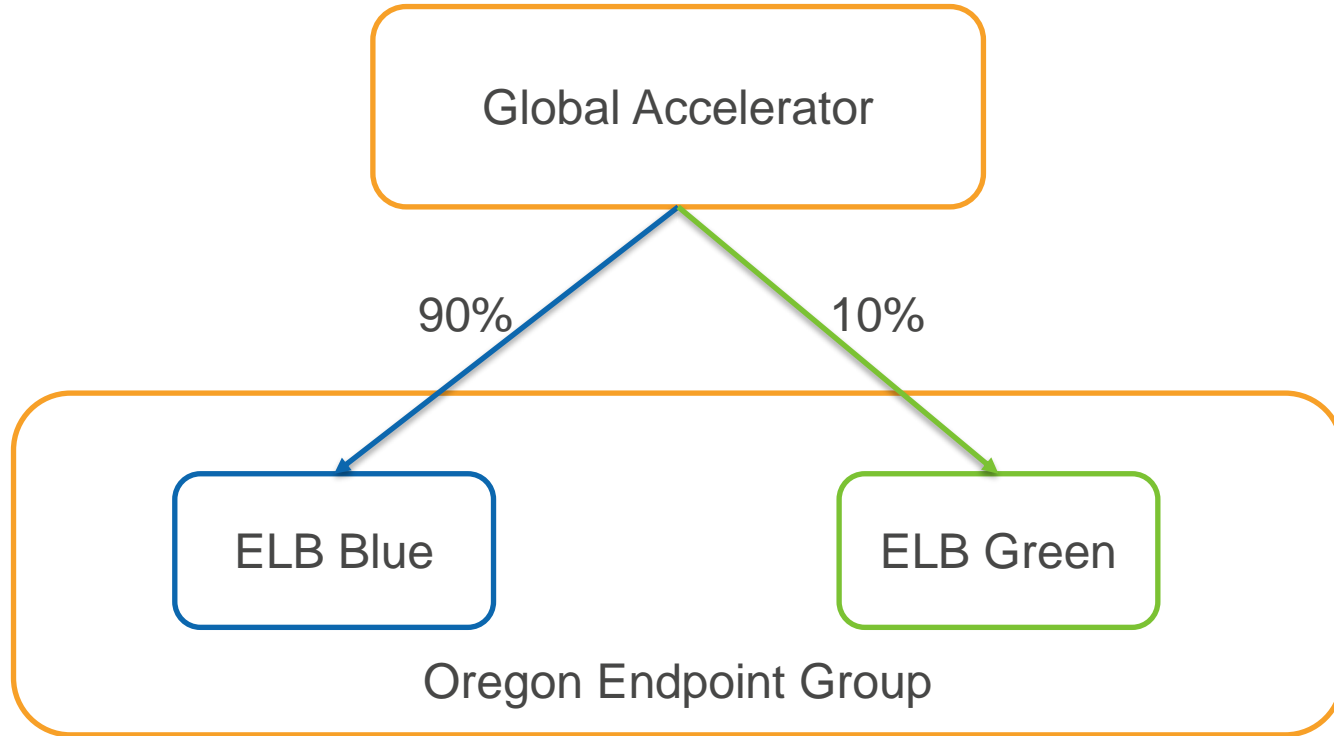


All traffic routed to Ohio

Traffic Dial – Restore traffic after deployment



Weight – Shape traffic flow inside a region



Summary

Reacts instantaneously to changes in traffic and health

Request routed through Edge network and AWS global network

Add regions, move regions without impact to client

DDoS protection


Lab – Global Accelerator

Setup ELB in multiple regions

Setup Global Accelerator

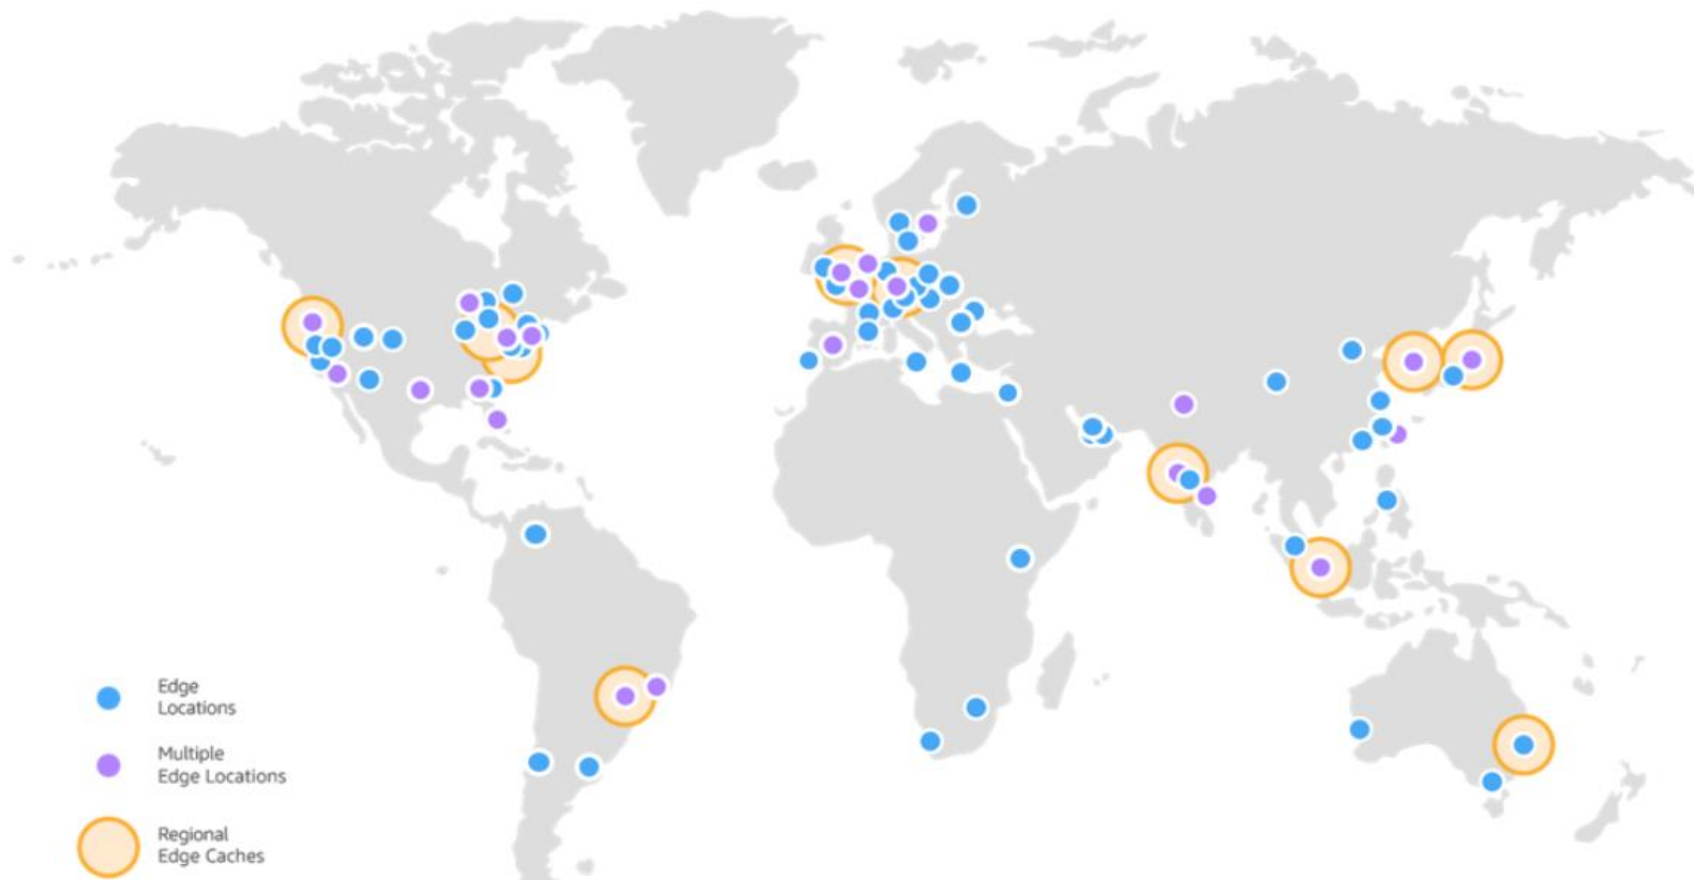
Adjust traffic dial to reroute traffic

Simulate Failure



Edge Locations: <https://aws.amazon.com/cloudfront/features/>

Locations and 11 Regional Edge Caches) in 84 cities across 42 countries. Amazon CloudFront Edge locations are located in:





Chandra Lingam

57,000+ Students



For AWS self-paced video courses, visit:

<https://www.cloudwavetraining.com/>

