



AWS
re:Invent

CON326-R

Running Kubernetes applications on AWS Fargate

Nathan Taber

Senior Product Manager
Amazon Web Services

Massimo Re Ferrè

Principal Developer Advocate
Amazon Web Services


**Make AWS the BEST PLACE
to run KUBERNETES**



**Production
Workloads**



**Native and
upstream**



**OSS
Contribution**



**Seamless
integrations**

Production workloads



Single tenant

Multi-AZ and highly available
architecture
by default

99.9% Service Level Agreement
for every cluster

Native and upstream



Upstream conformant

Integration testing
with **Kubernetes tooling**

APIs and existing tooling
just work

OSS contributions



AWS contributes
bug fixes, security patches, and
tooling improvements

Open-source components
Contribute to or maintain over
30 OSS projects on GitHub for
Kubernetes

Seamless integrations



Identity

Audits

Routing

Compliance

Monitoring

Logging

Ingress


Security

Databases

The background of the slide is a dark charcoal gray. It is decorated with numerous isometric wireframe cubes of various sizes. These cubes are drawn with thin, light blue lines. They are scattered across the entire frame, with some appearing in the foreground and others receding into the background, creating a sense of depth. The cubes are not solid; they are just outlines, allowing the dark background to show through them.

All the building blocks for Kubernetes
in one place

Containers options on AWS – over time




Containers options on AWS – over time




AWS Cloud

Docker
Host



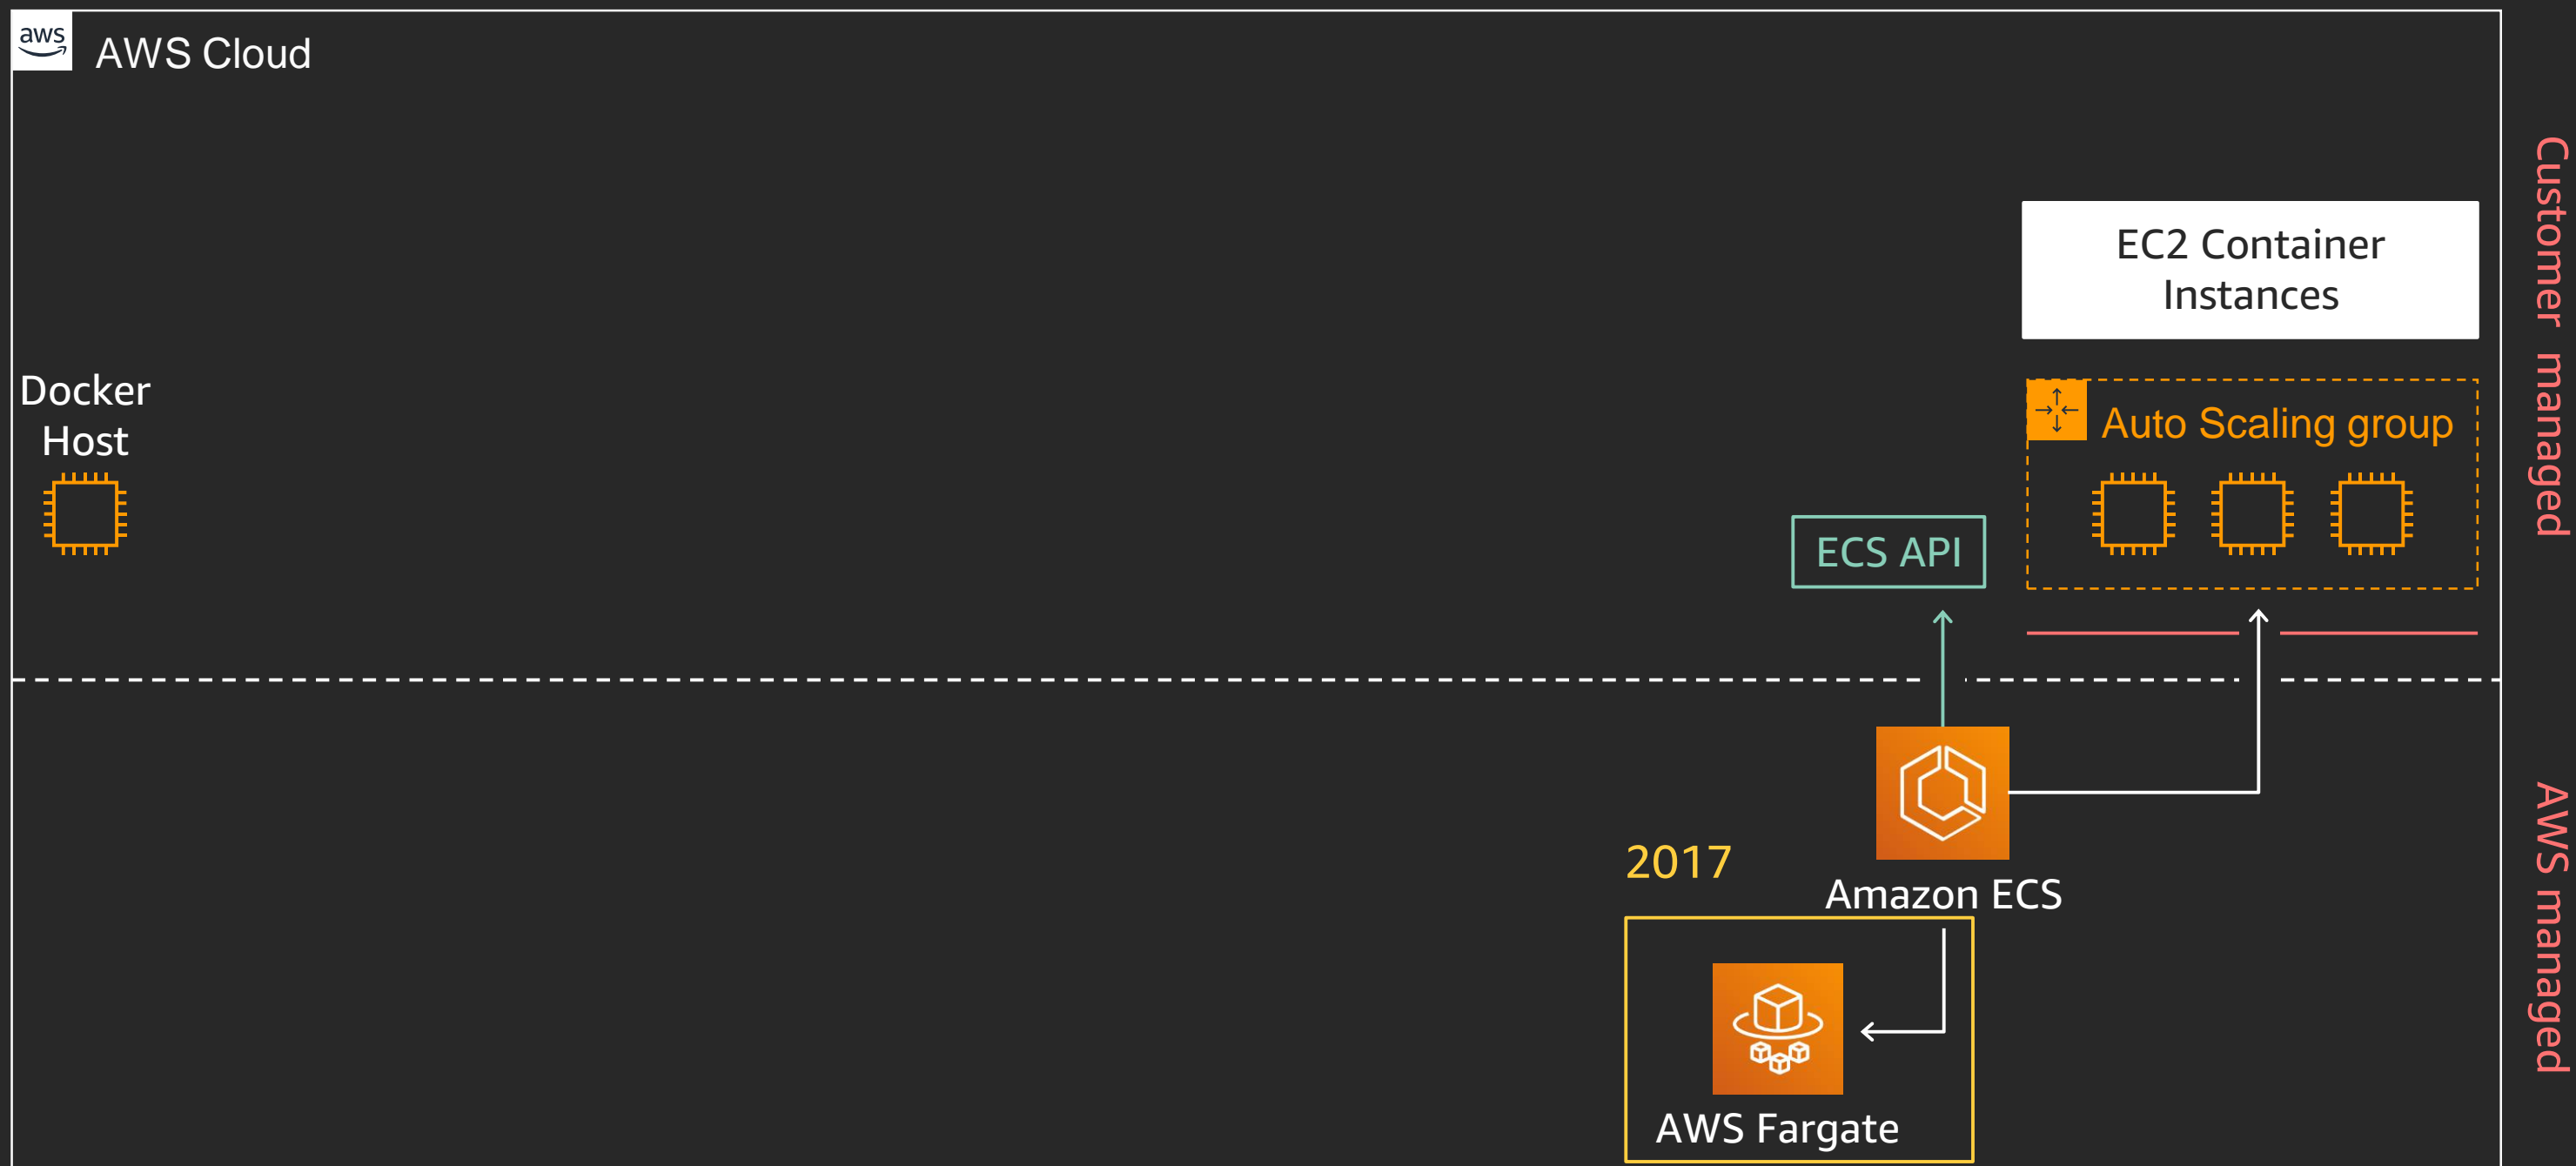
2015



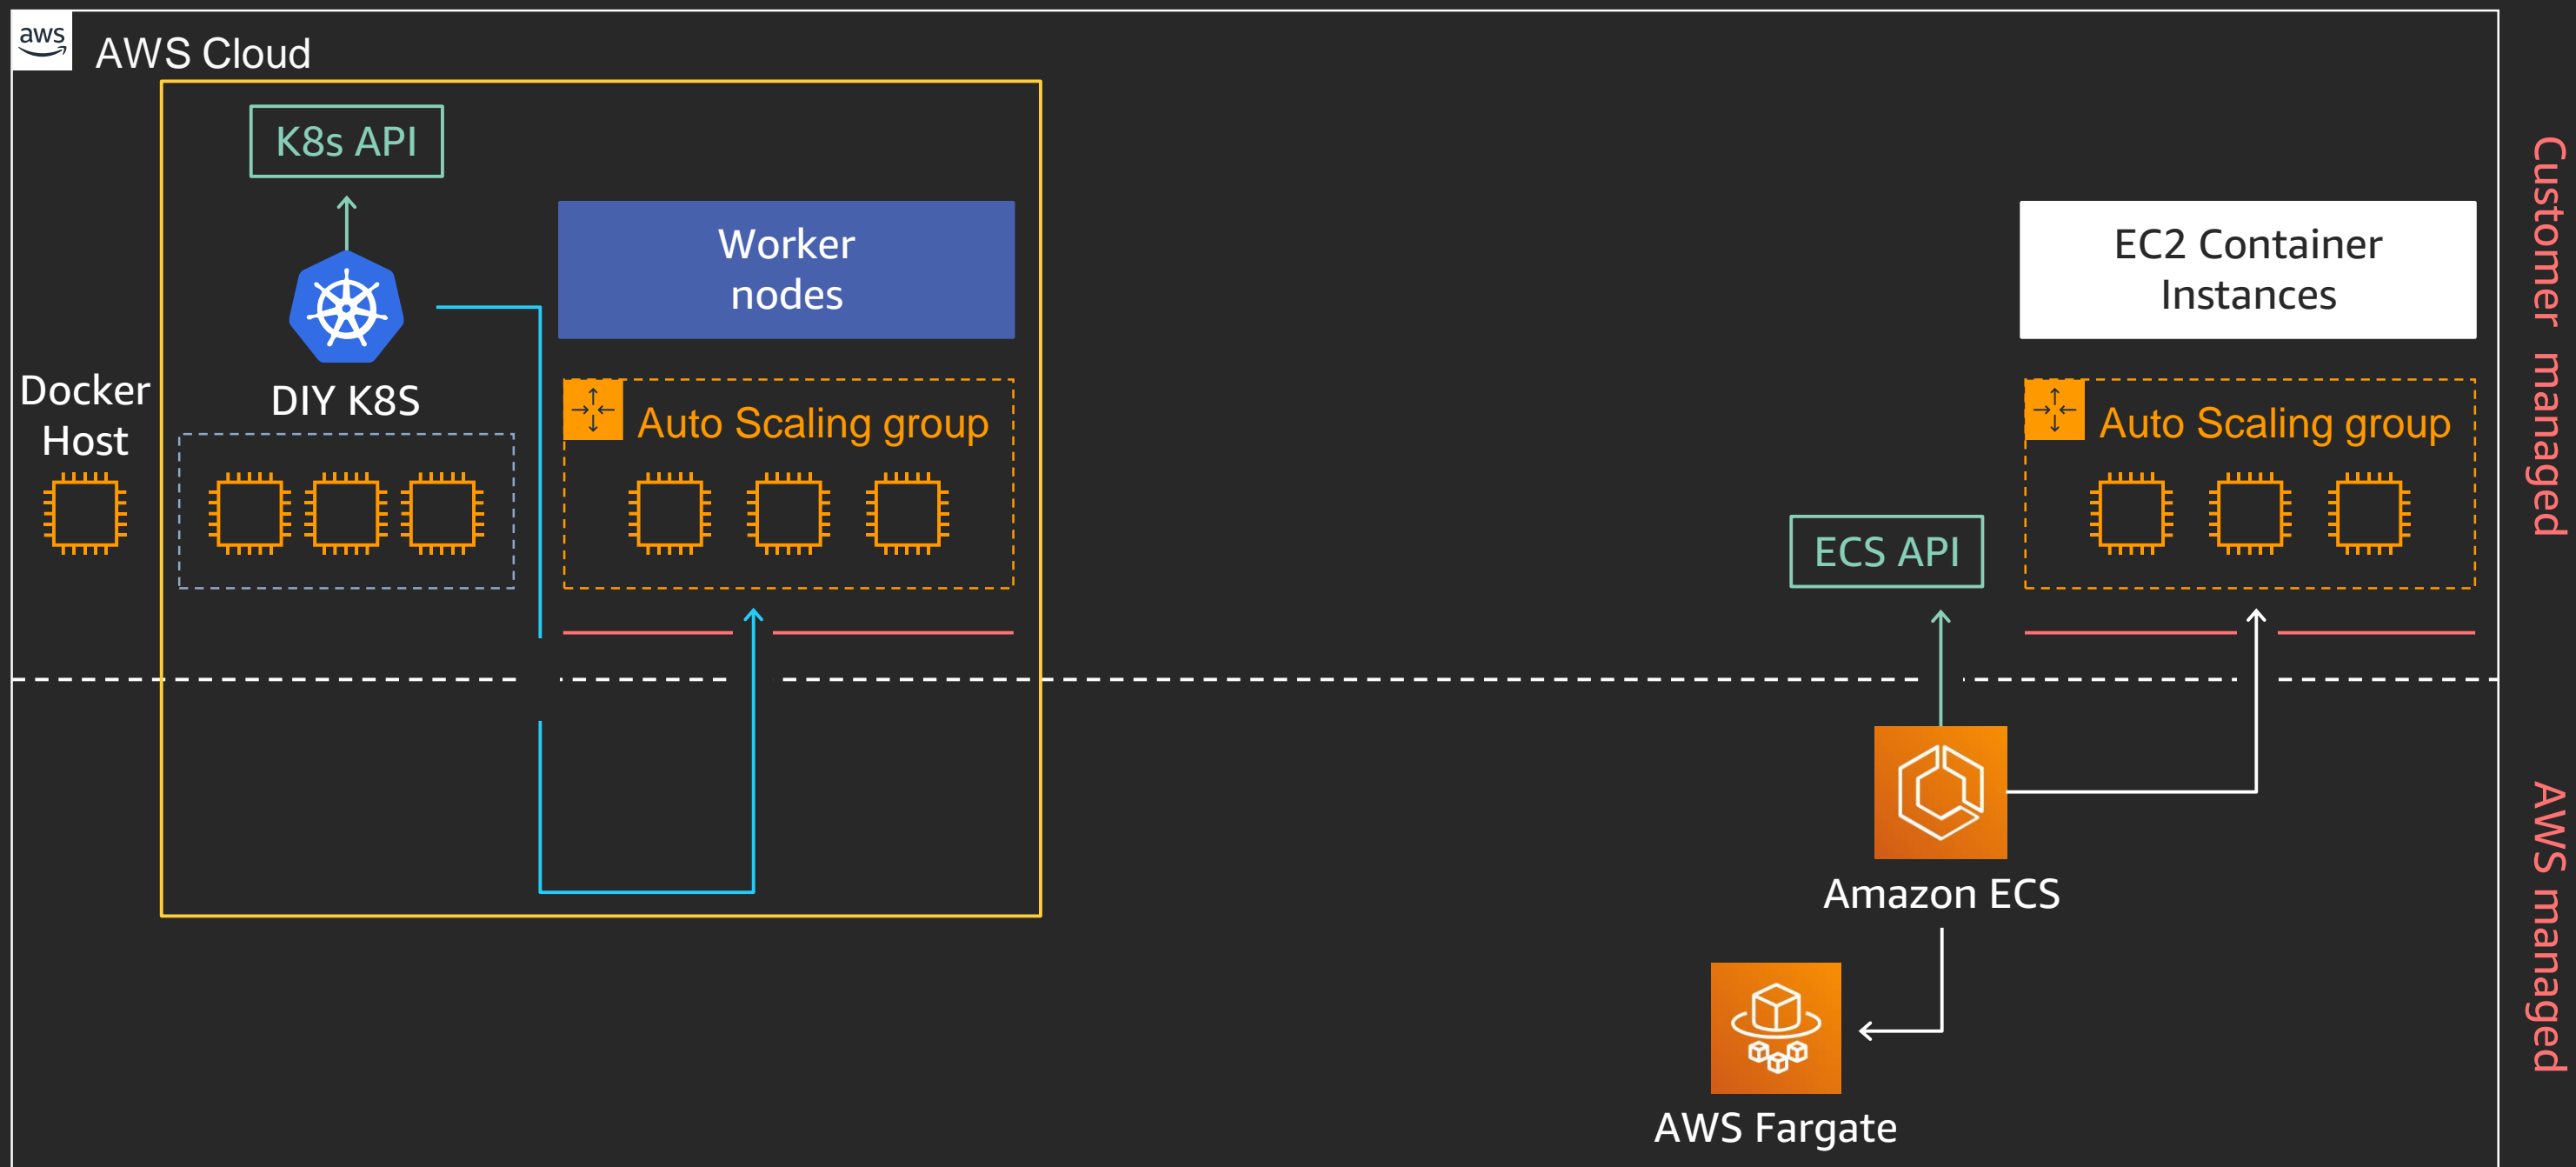
Customer managed

AWS managed

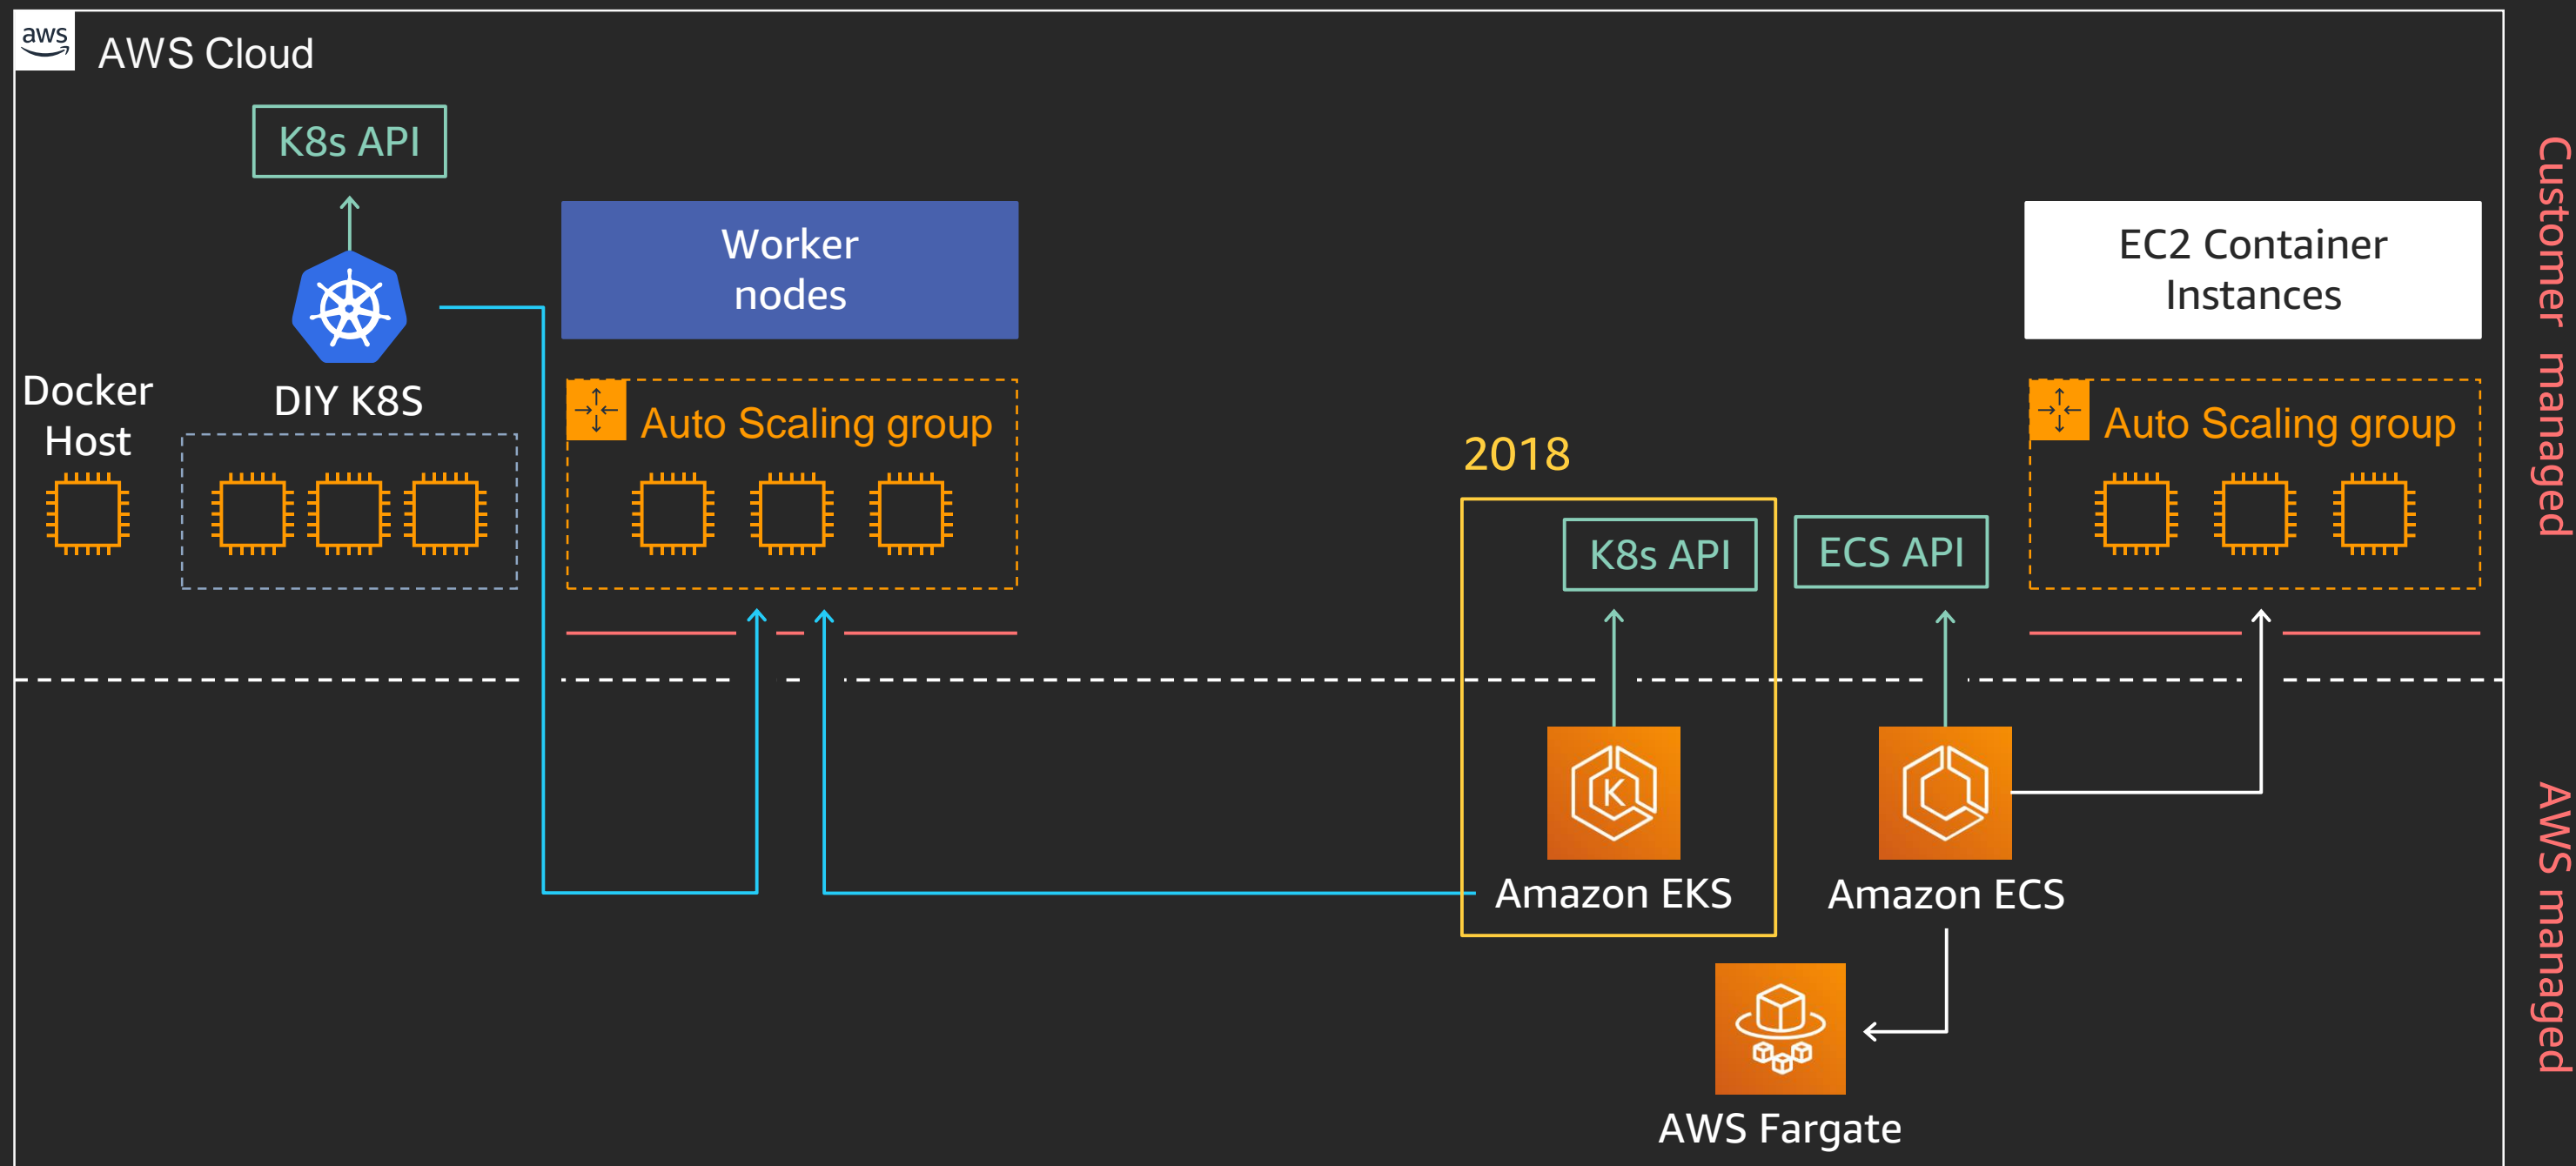
Containers options on AWS – over time



Containers options on AWS – over time




Containers options on AWS – over time



Phase 1

Management of the Kubernetes control plane




Phase 1

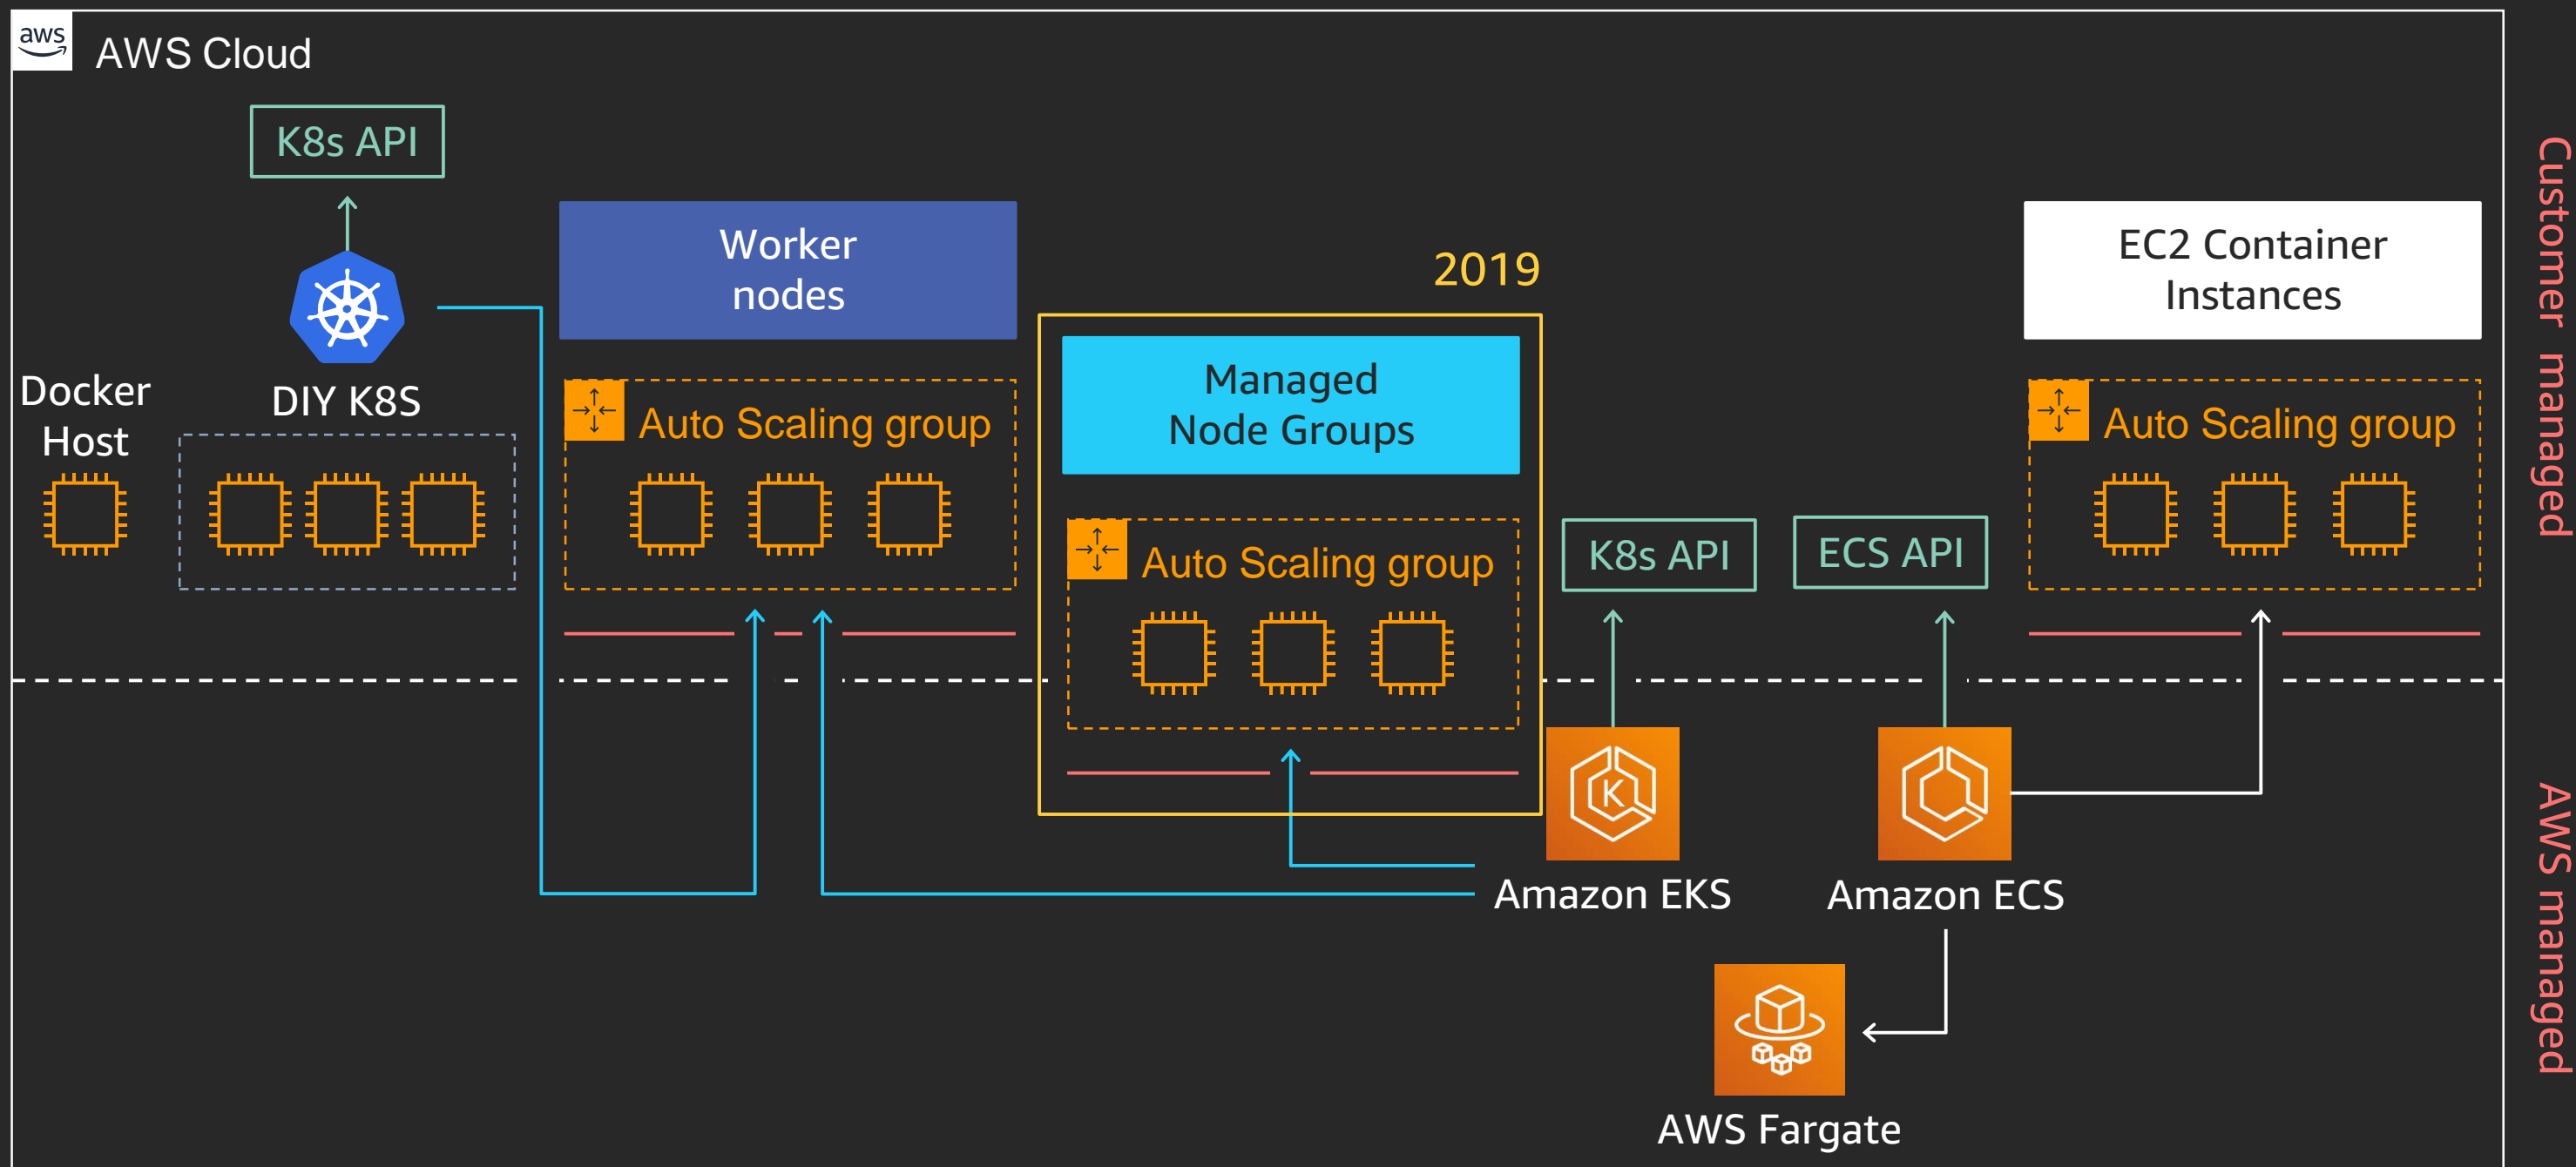
Management of the
Kubernetes control plane

Phase 2

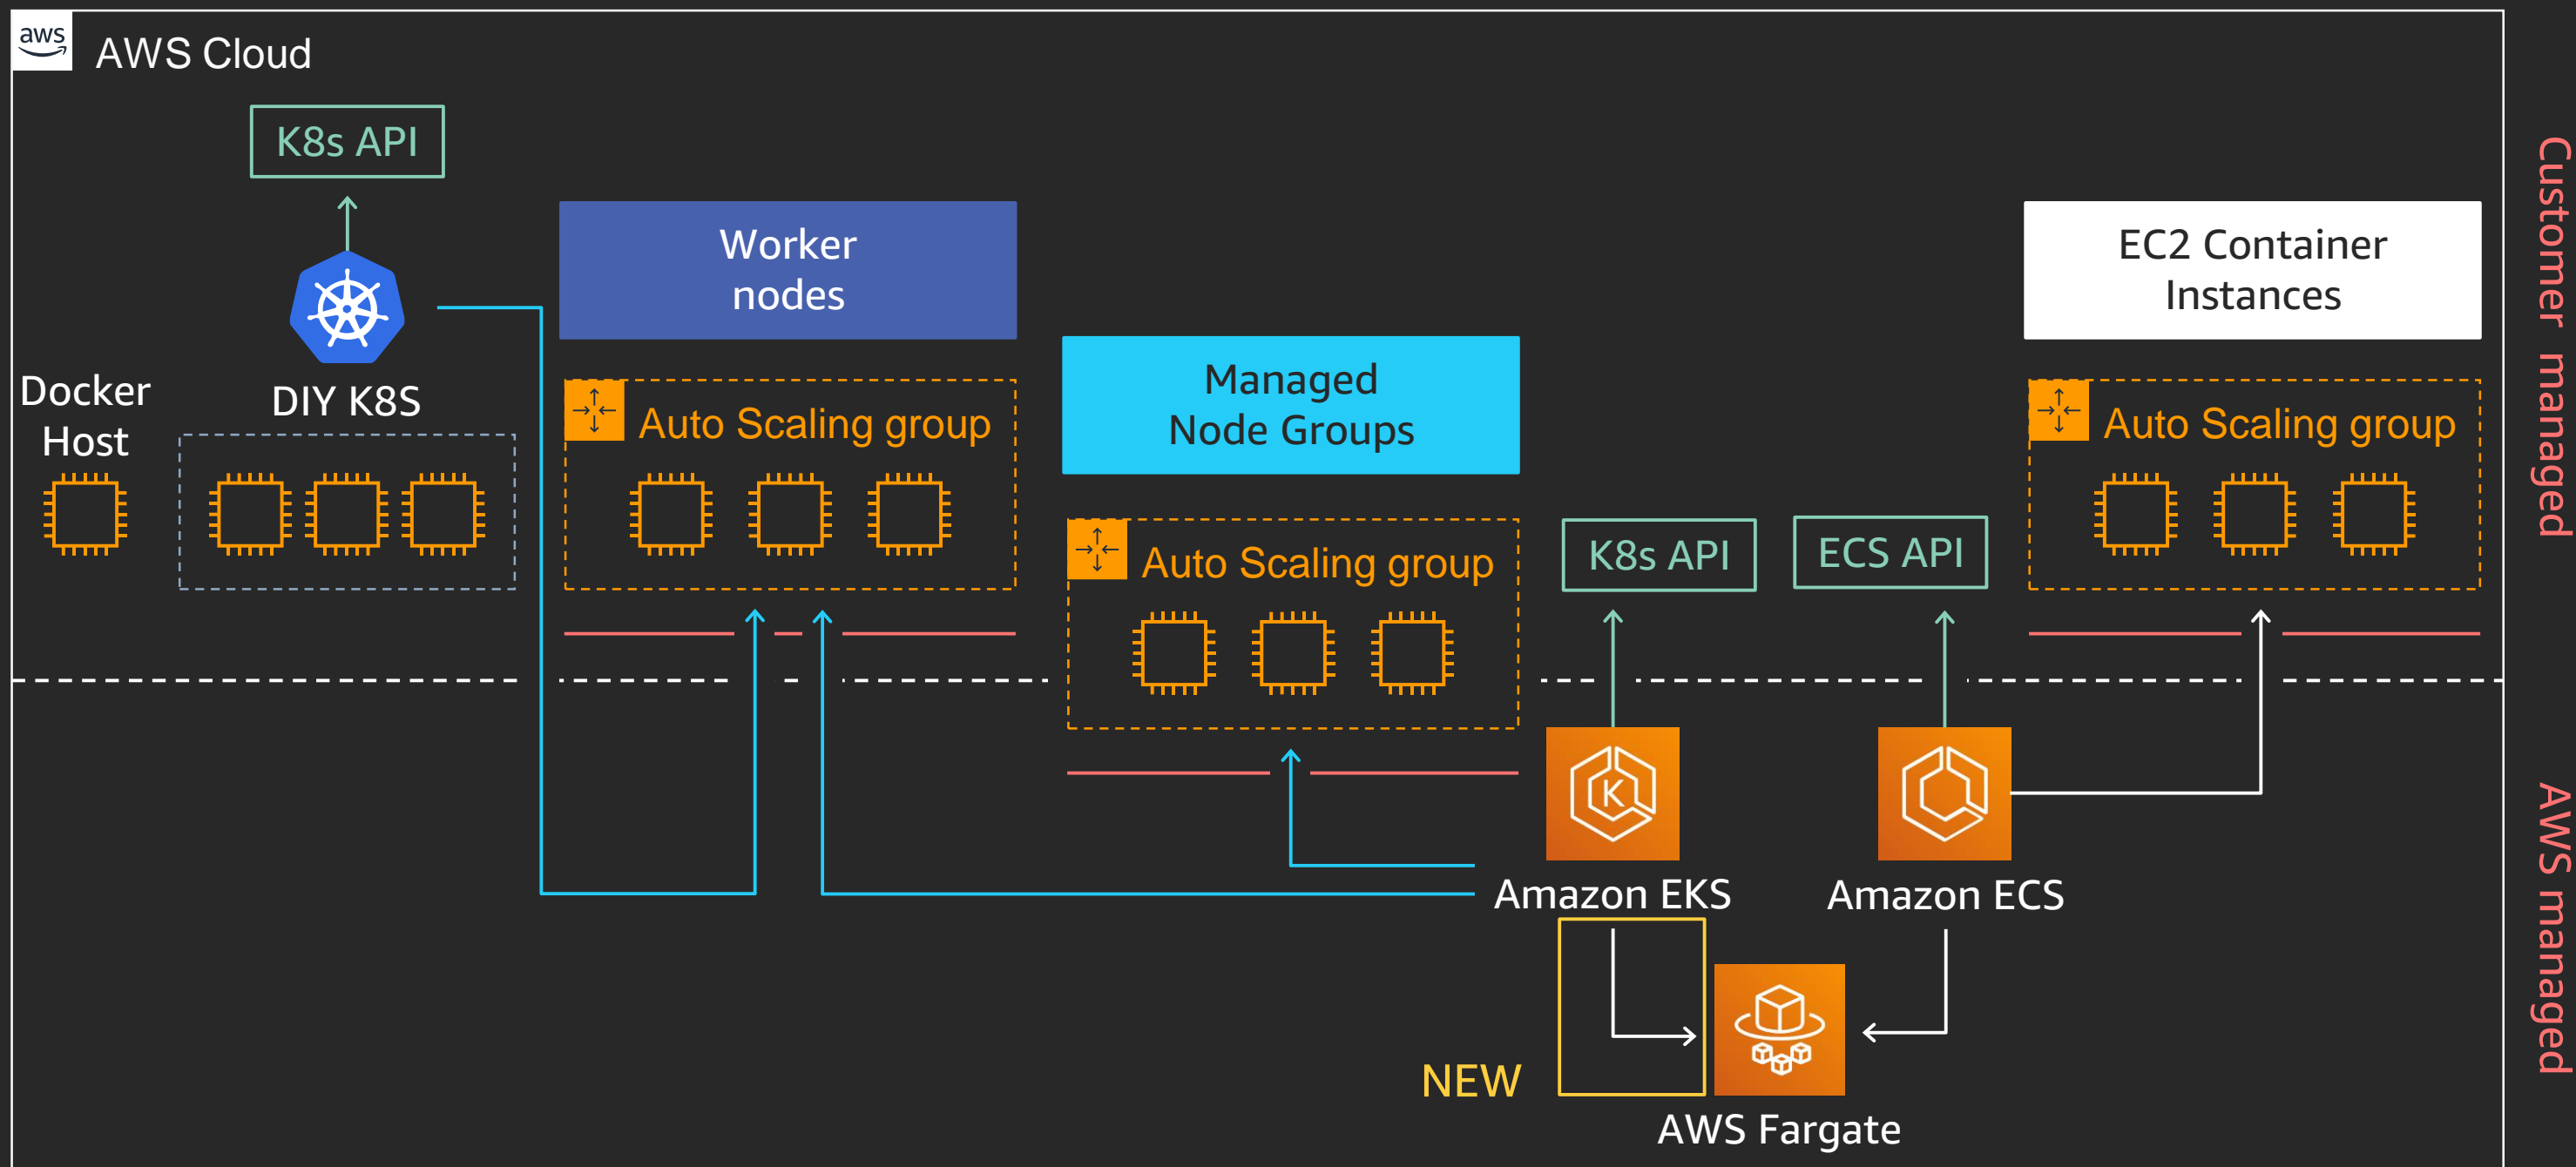
**Management of the
Kubernetes data plane**



Containers options on AWS – over time



Containers options on AWS – over time




You should be able to write your code and have it run, without having to worry about configuring complex management tools.

This is the vision behind AWS Fargate.

Dr. Werner Vogels

CTO, Amazon.com

Amazon EKS on Fargate



Bring existing pods

You don't need to change your existing pods.


Fargate works with existing workflows and services that run on Kubernetes.



Production ready

Launch pods quickly. Easily run pods across multiple AZs for high availability.

Each pod runs in an isolated compute environment.



Rightsized and integrated


Only pay for the resources you need to run your pods.

Includes native AWS integrations for networking and security.

What matters for Fargate



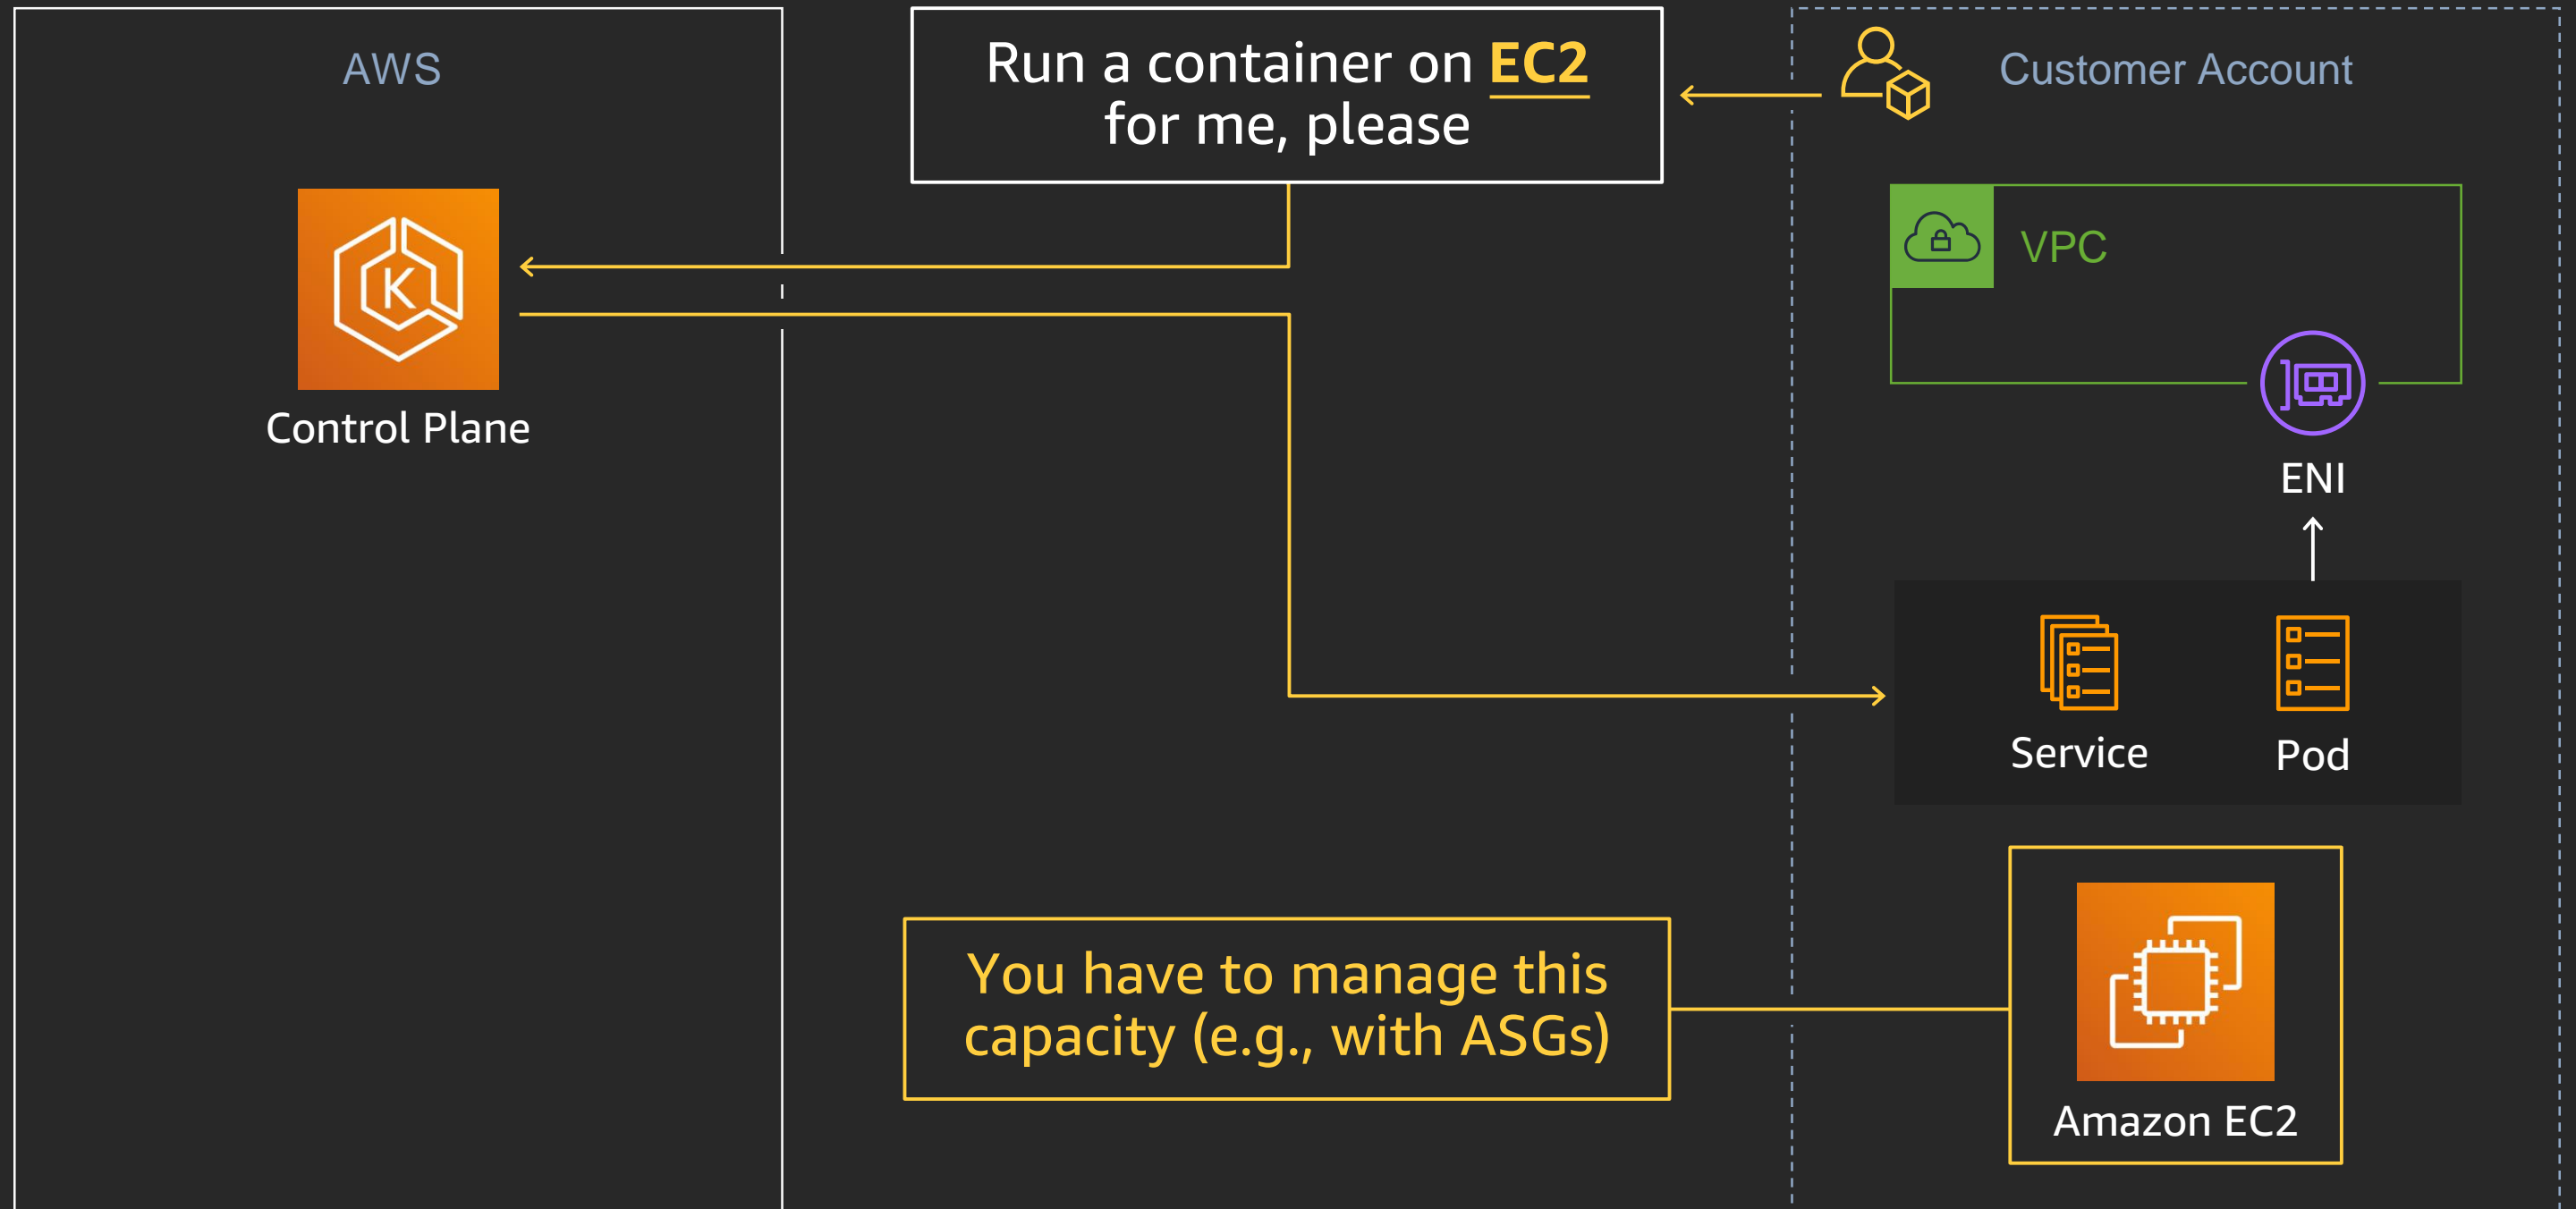
Fargate is a serverless compute platform for containers on AWS



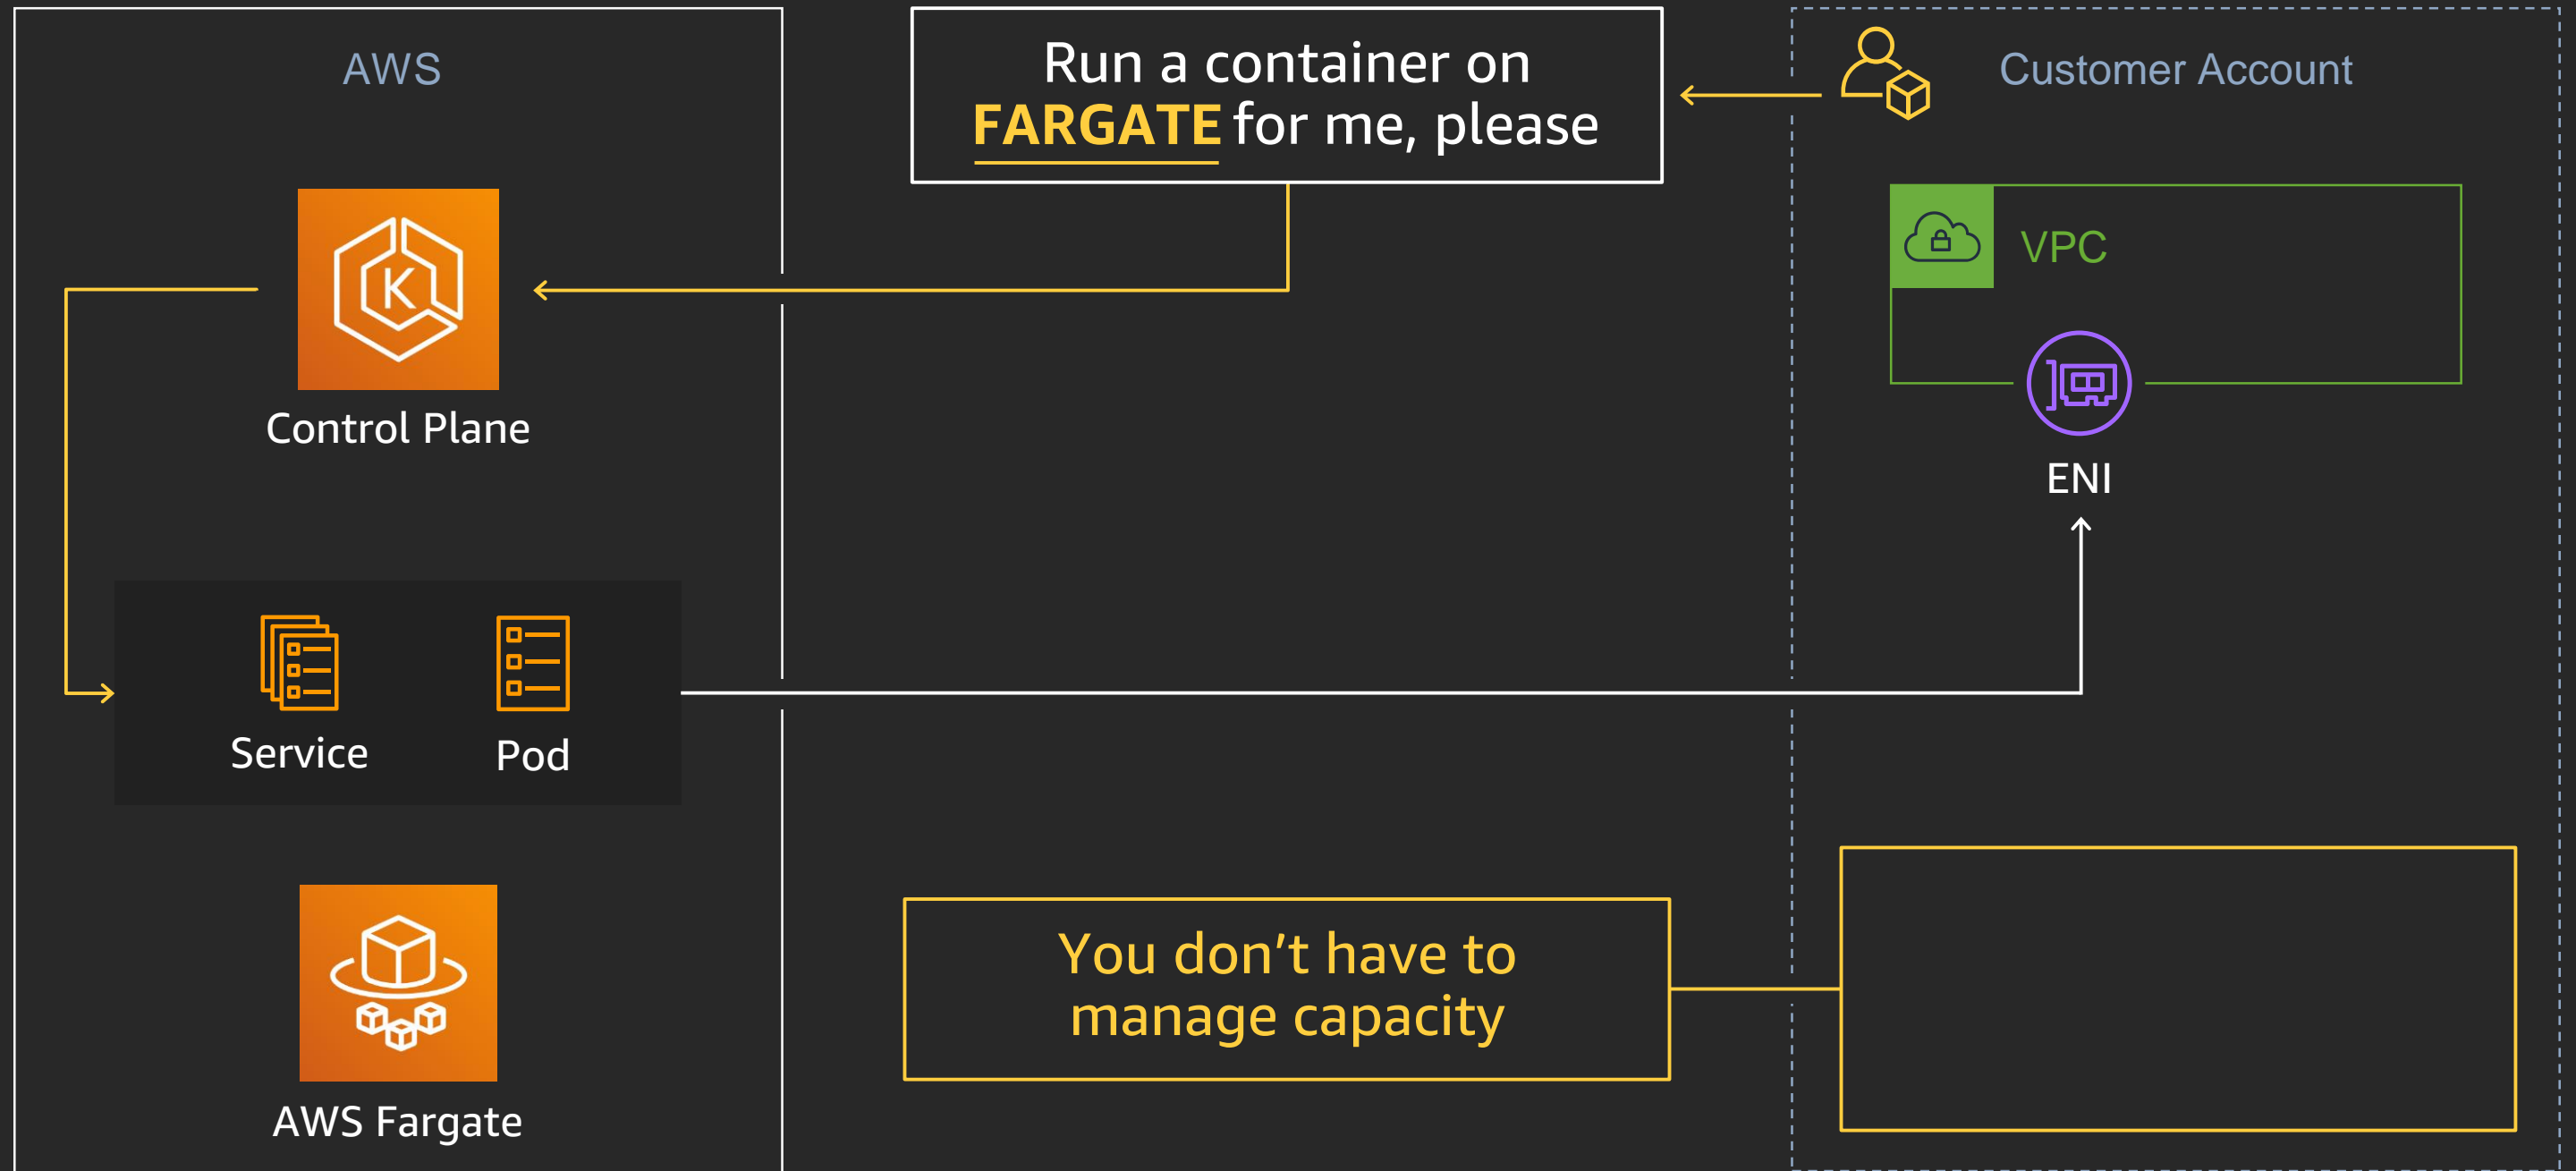
The differences between using EKS and ECS with Fargate are driven by the orchestration system

How EKS works on Fargate

The EC2 flow at 33,000 feet



The Fargate flow at 33,000 feet



Fargate Vs. (Un)Managed Nodes

	Fargate	Managed nodes	Unmanaged nodes
Units of work	Pod	Pod and EC2	Pod and EC2
Unit of charge	Pod	EC2	EC2

Fargate Vs. (Un)Managed Nodes

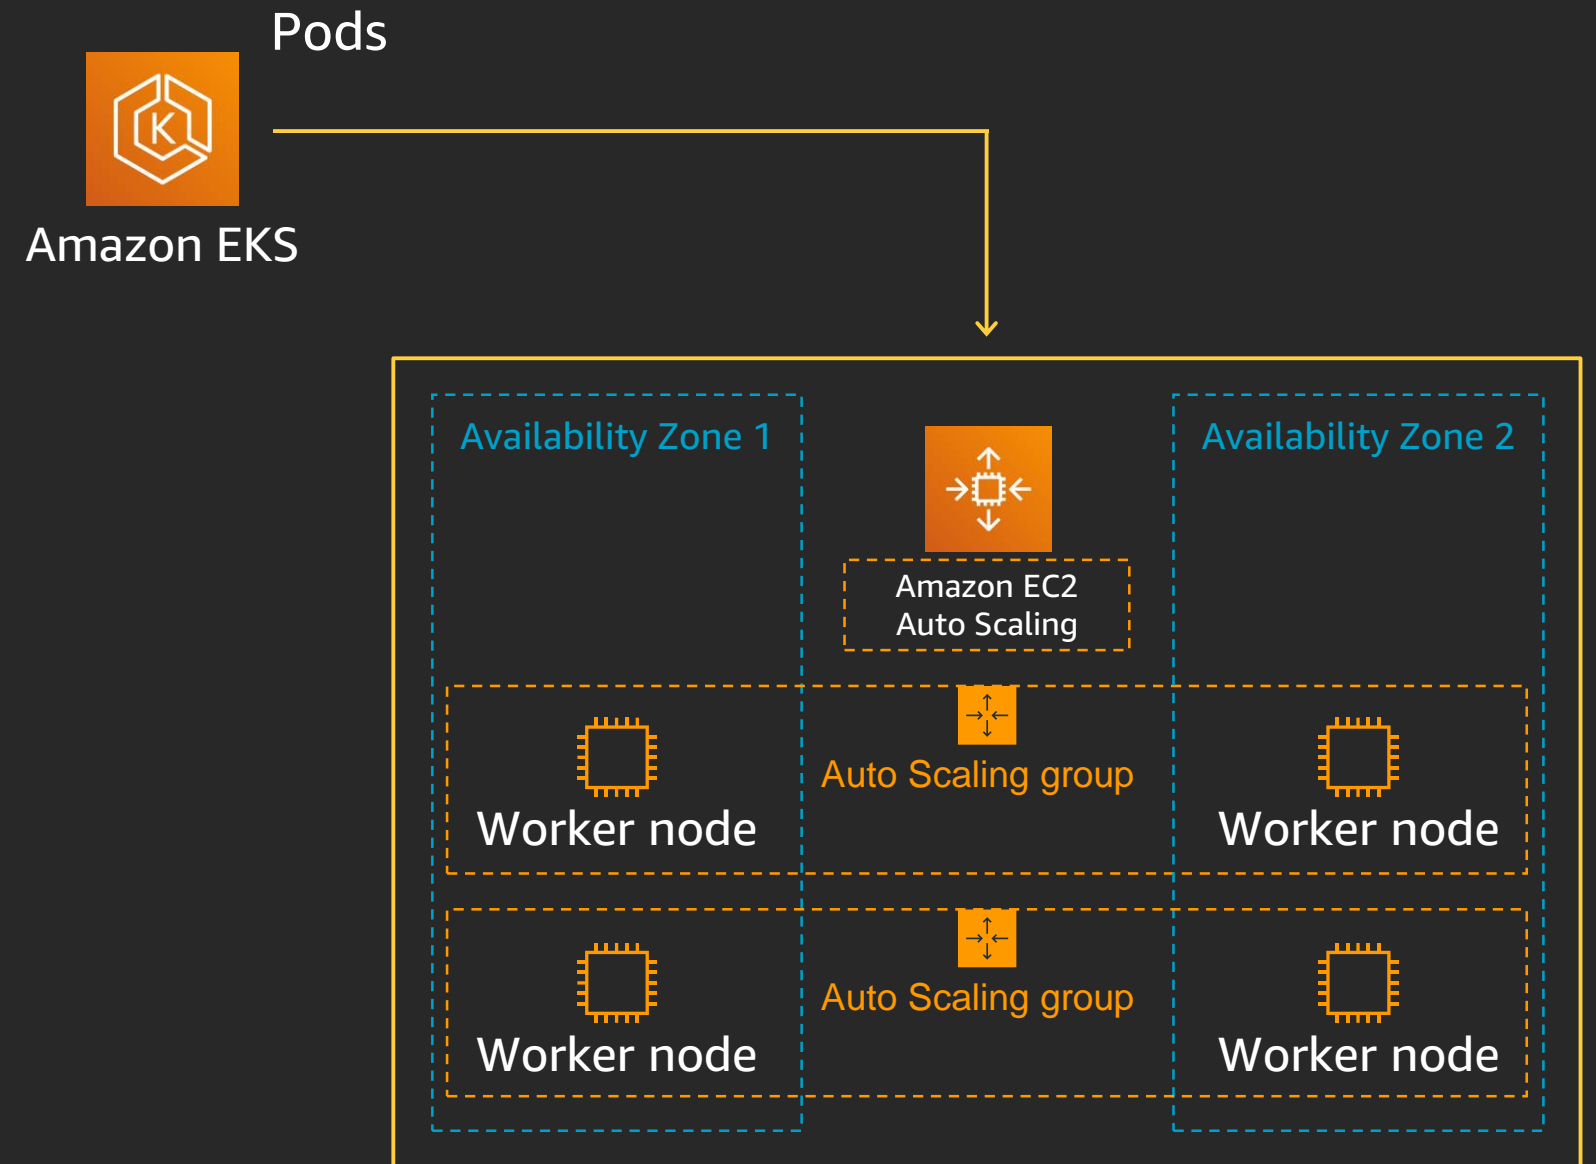
	Fargate	Managed nodes	Unmanaged nodes
Units of work	Pod	Pod and EC2	Pod and EC2
Unit of charge	Pod	EC2	EC2
Host lifecycle	There is no visible host	AWS (SSH is allowed)	Customer
Host AMI	There is no visible host	AWS vetted AMIs	Customer BYO

Fargate vs. (Un)Managed Nodes

	Fargate	Managed nodes	Unmanaged nodes
Units of work	Pod	Pod and EC2	Pod and EC2
Unit of charge	Pod	EC2	EC2
Host lifecycle	There is no visible host	AWS (SSH is allowed)	Customer
Host AMI	There is no visible host	AWS vetted AMIs	Customer BYO
Host : Pods	1 : 1	1 : many	1 : many

EKS data plane options

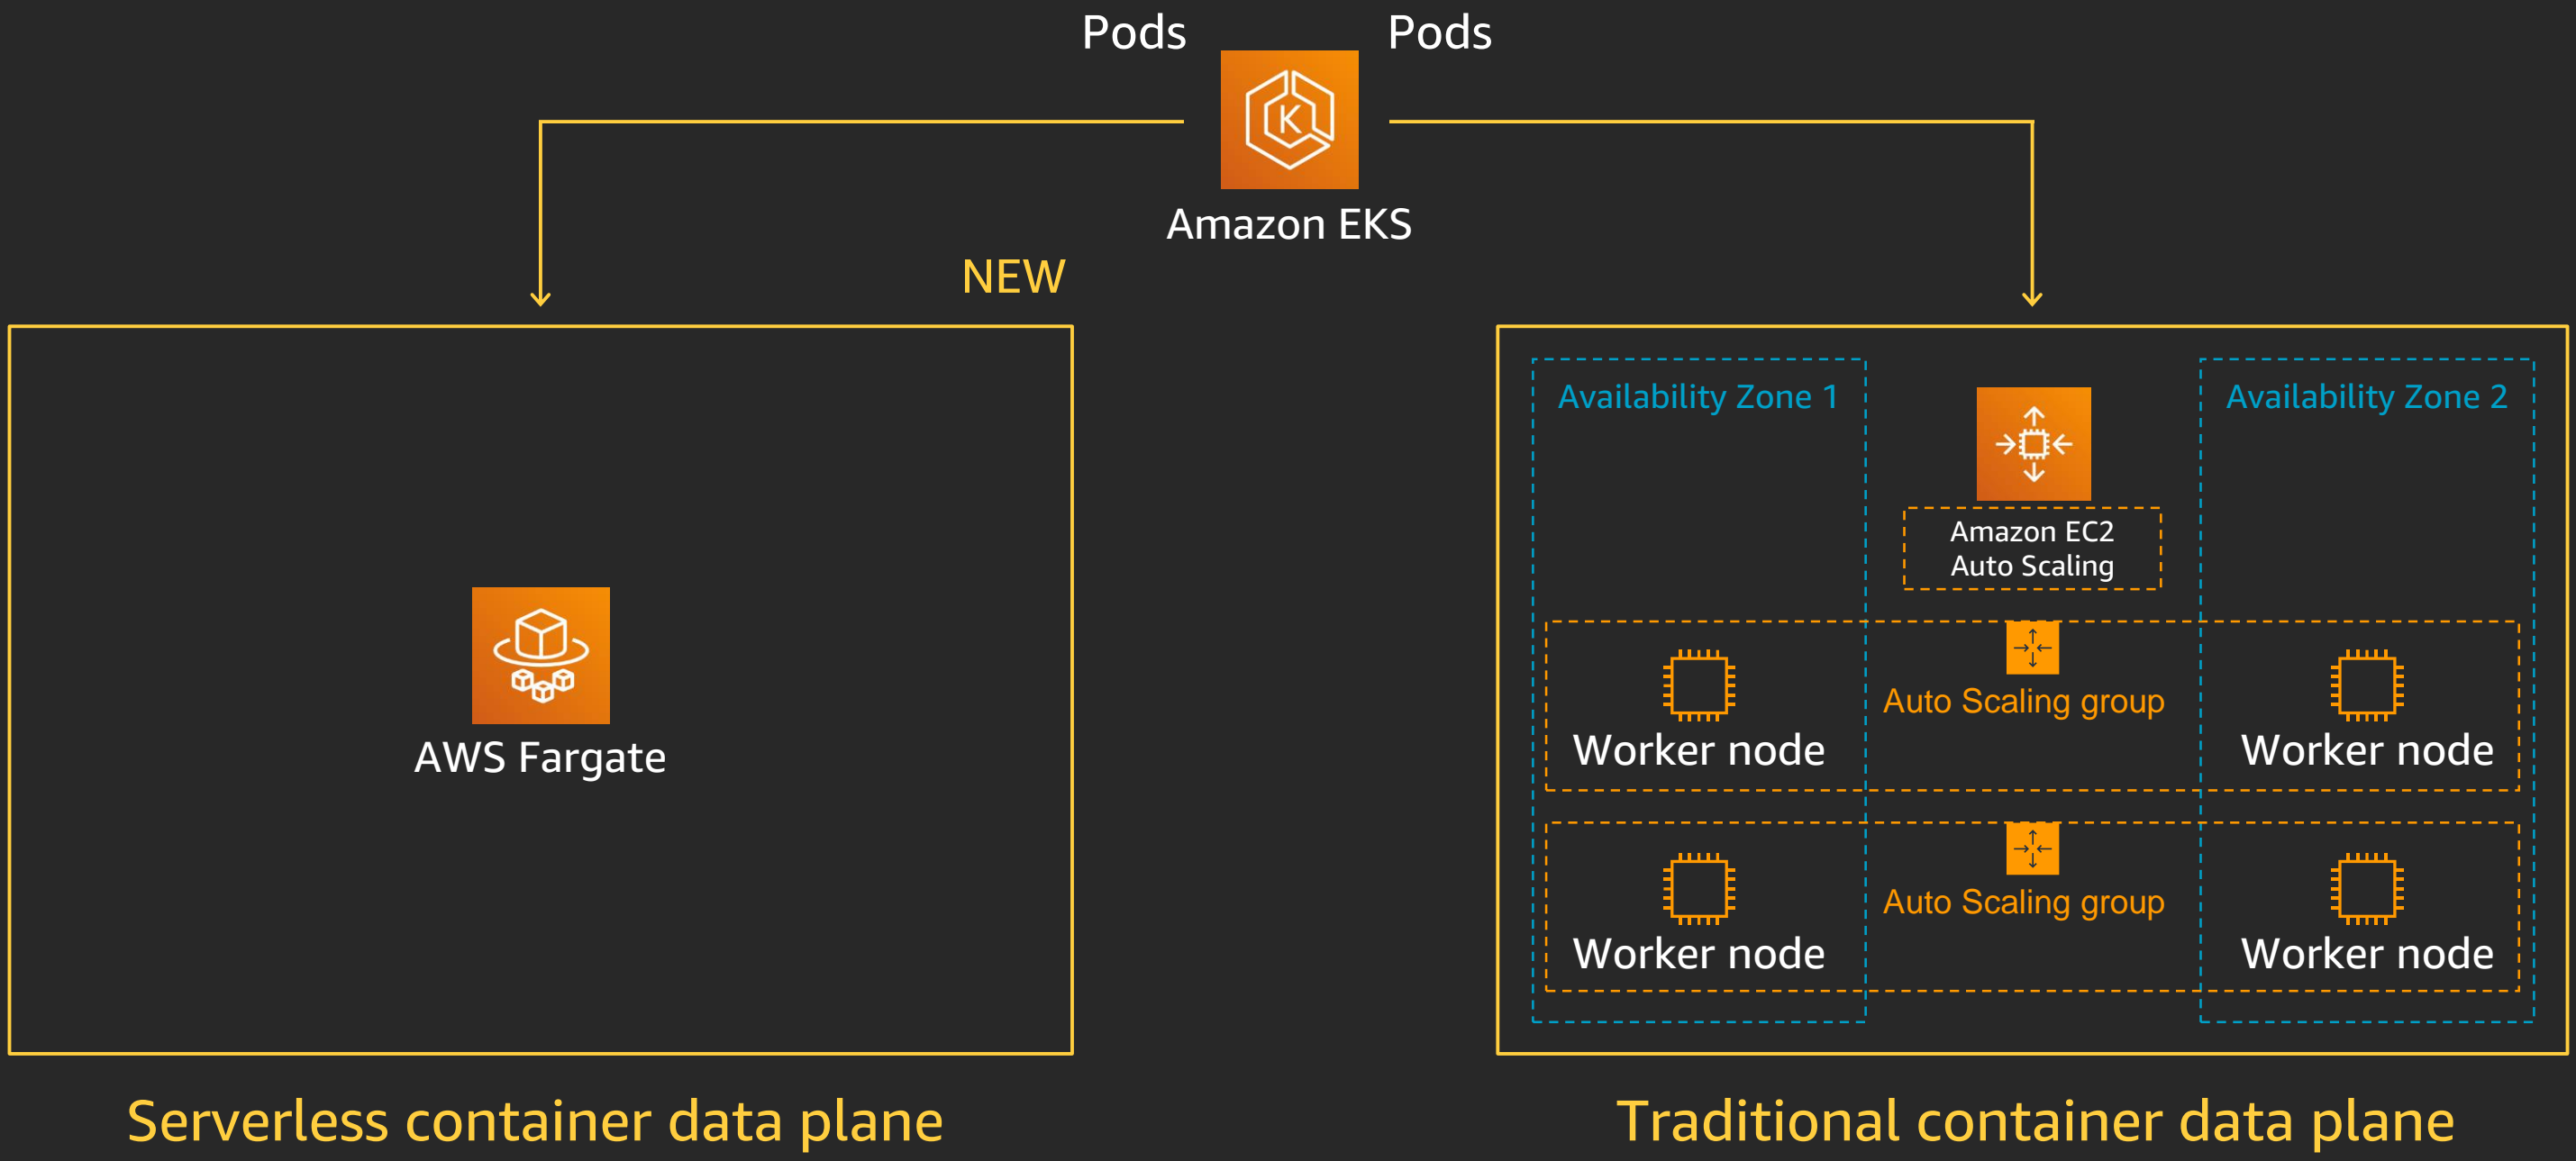
Worker nodes only



Traditional container data plane

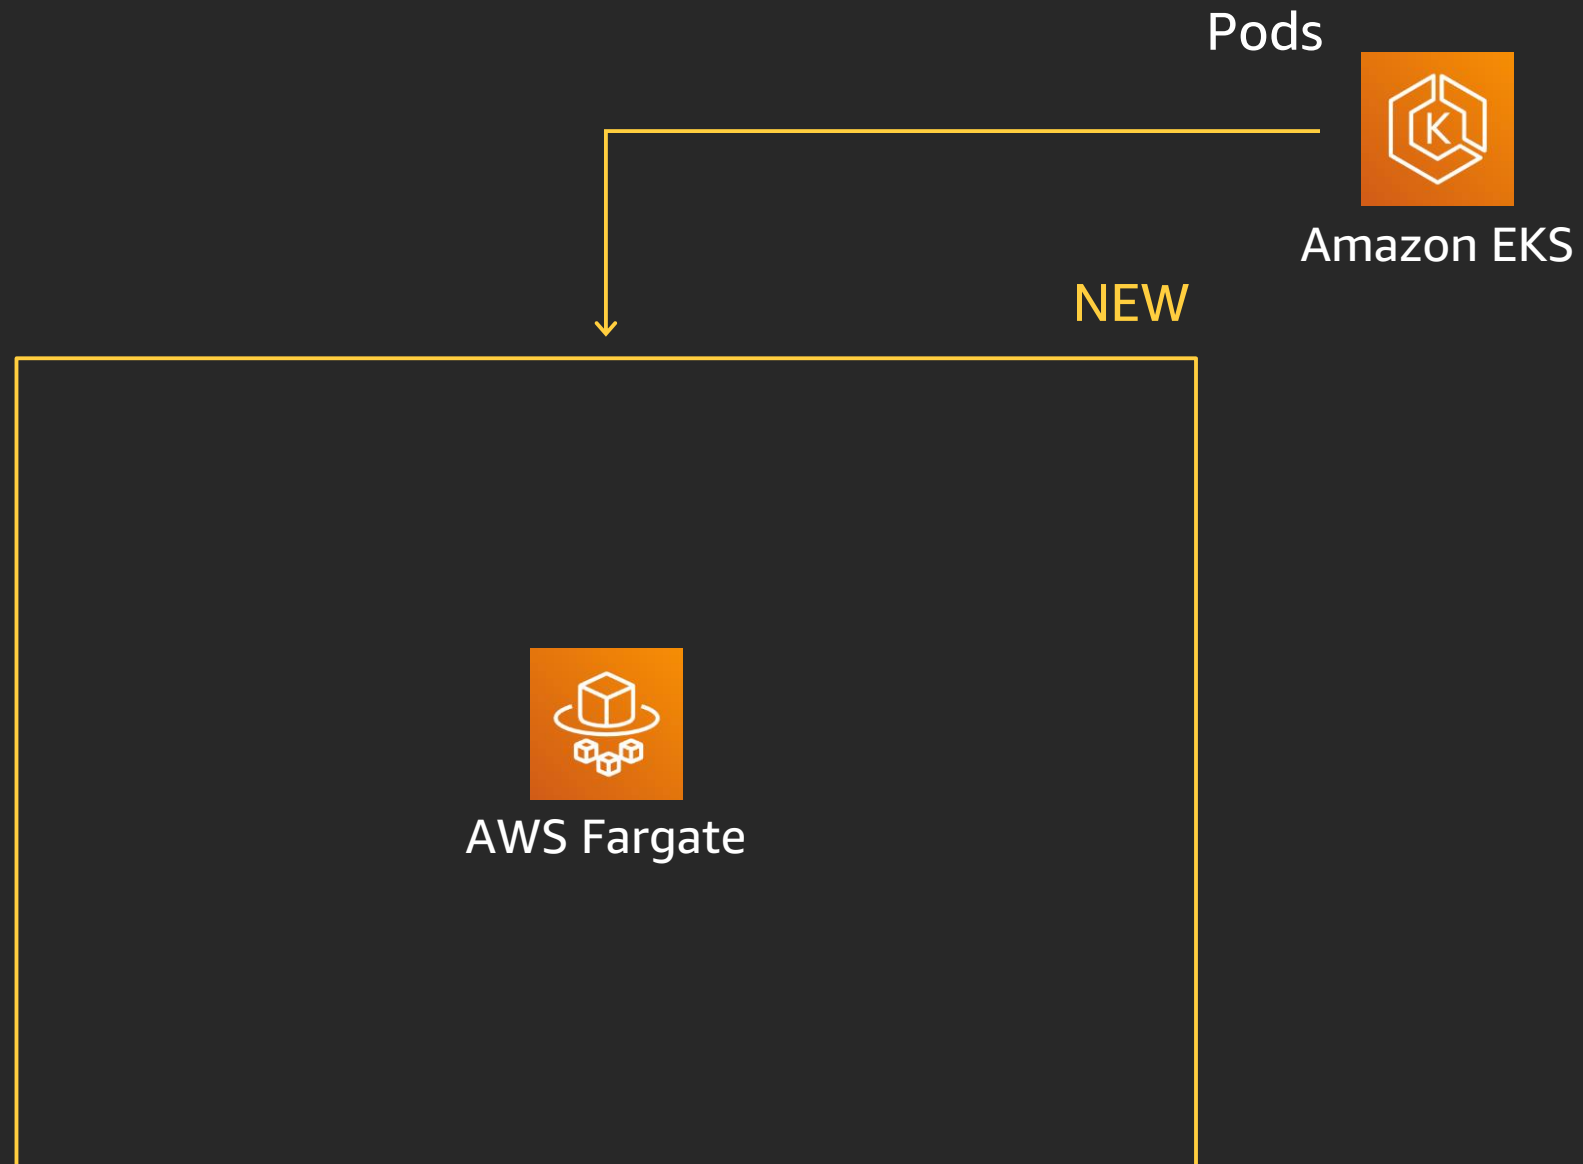
EKS data plane options

Mixed mode



EKS data plane options

Fargate only




Serverless container data plane

Demo

Kubernetes and EKS: Objects and constructs

Amazon EKS

Kubernetes and EKS: Objects and constructs



Fargate profile template

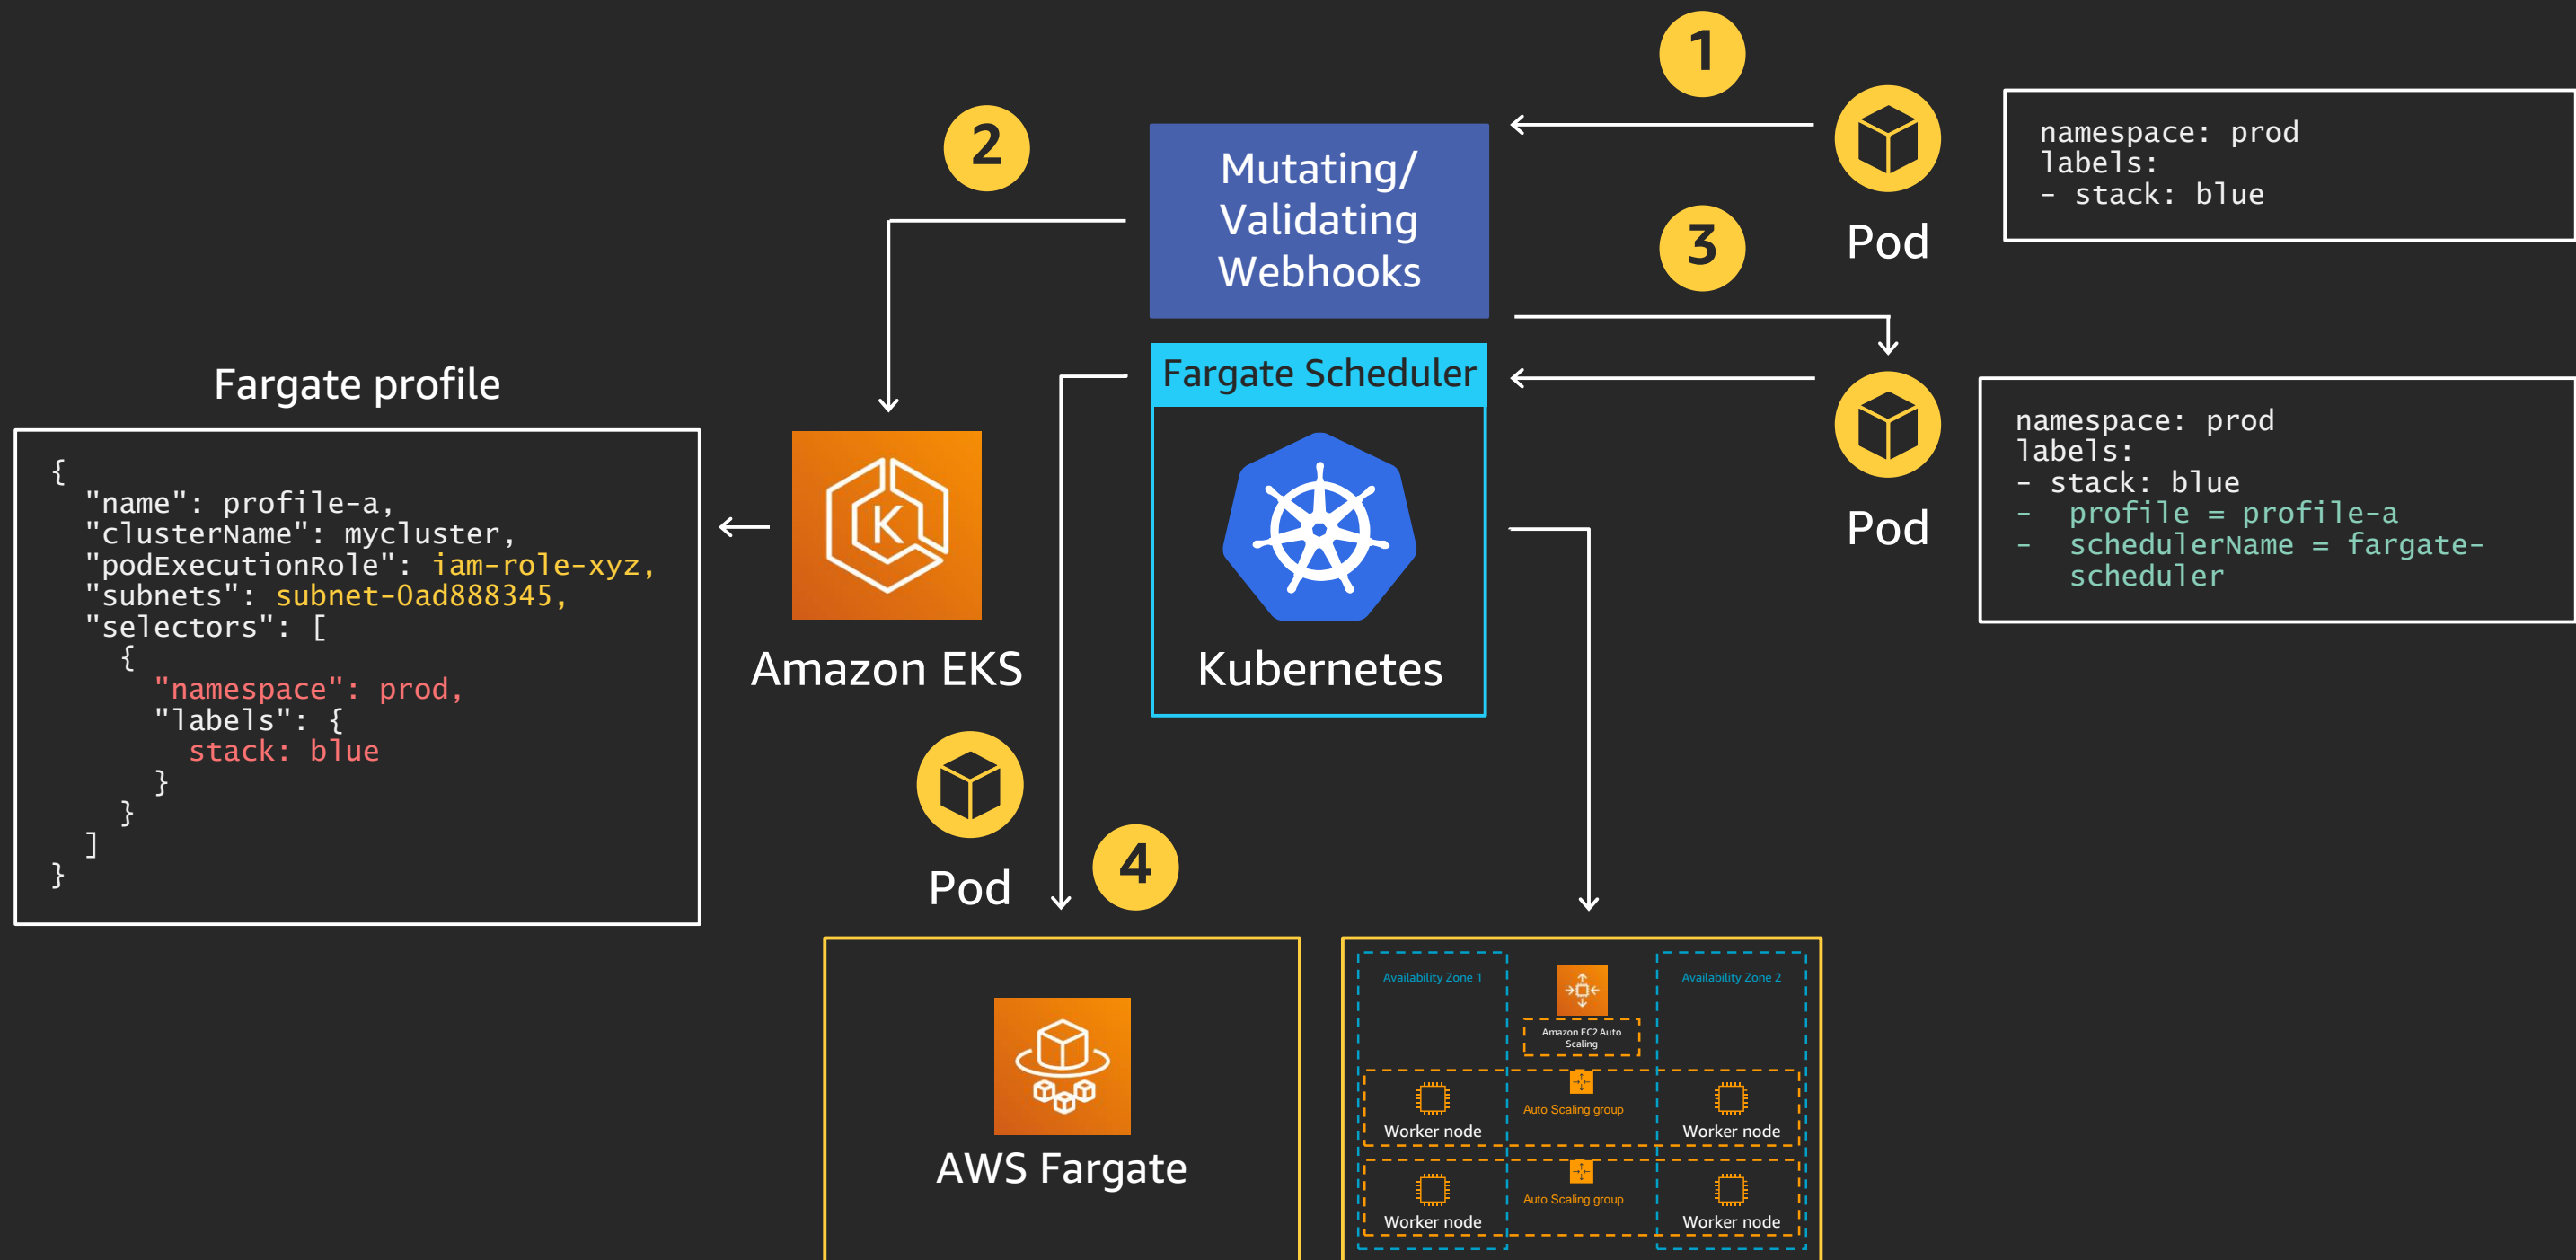
```
{
  "status": "ACTIVE",
  "subnets": [
    "subnet-0de8355bc4ds45af3",
    "subnet-0det555b36hdy67d3"
  ],
  "clusterName": "FargateCluster",
  "fargateProfileArn": "arn:aws:eks:us-west-2:123456789:fargateprofile/FargateCluster/FargateProfileCatchAll/4cg3303c-539e-a202-5b75-bb1dd3dd0590",
  "selectors": [
    {
      "namespace": "default"
    },
    {
      "namespace": "kube-system"
    },
    {
      "labels": {
        "foo": "bar"
      },
      "namespace": "mynamespace"
    }
  ],
  "fargateProfileName": "FargateProfileCatchAll",
  "podExecutionRole": "arn:aws:iam::123456789:role/FargateCluster-SERVICE-ROLE-AWSServiceRoleFargateCluster-1PLJY3220ID6I",
  "createdAt": 1573039680.227
}
```

Subnets to pick for the Pod deployment

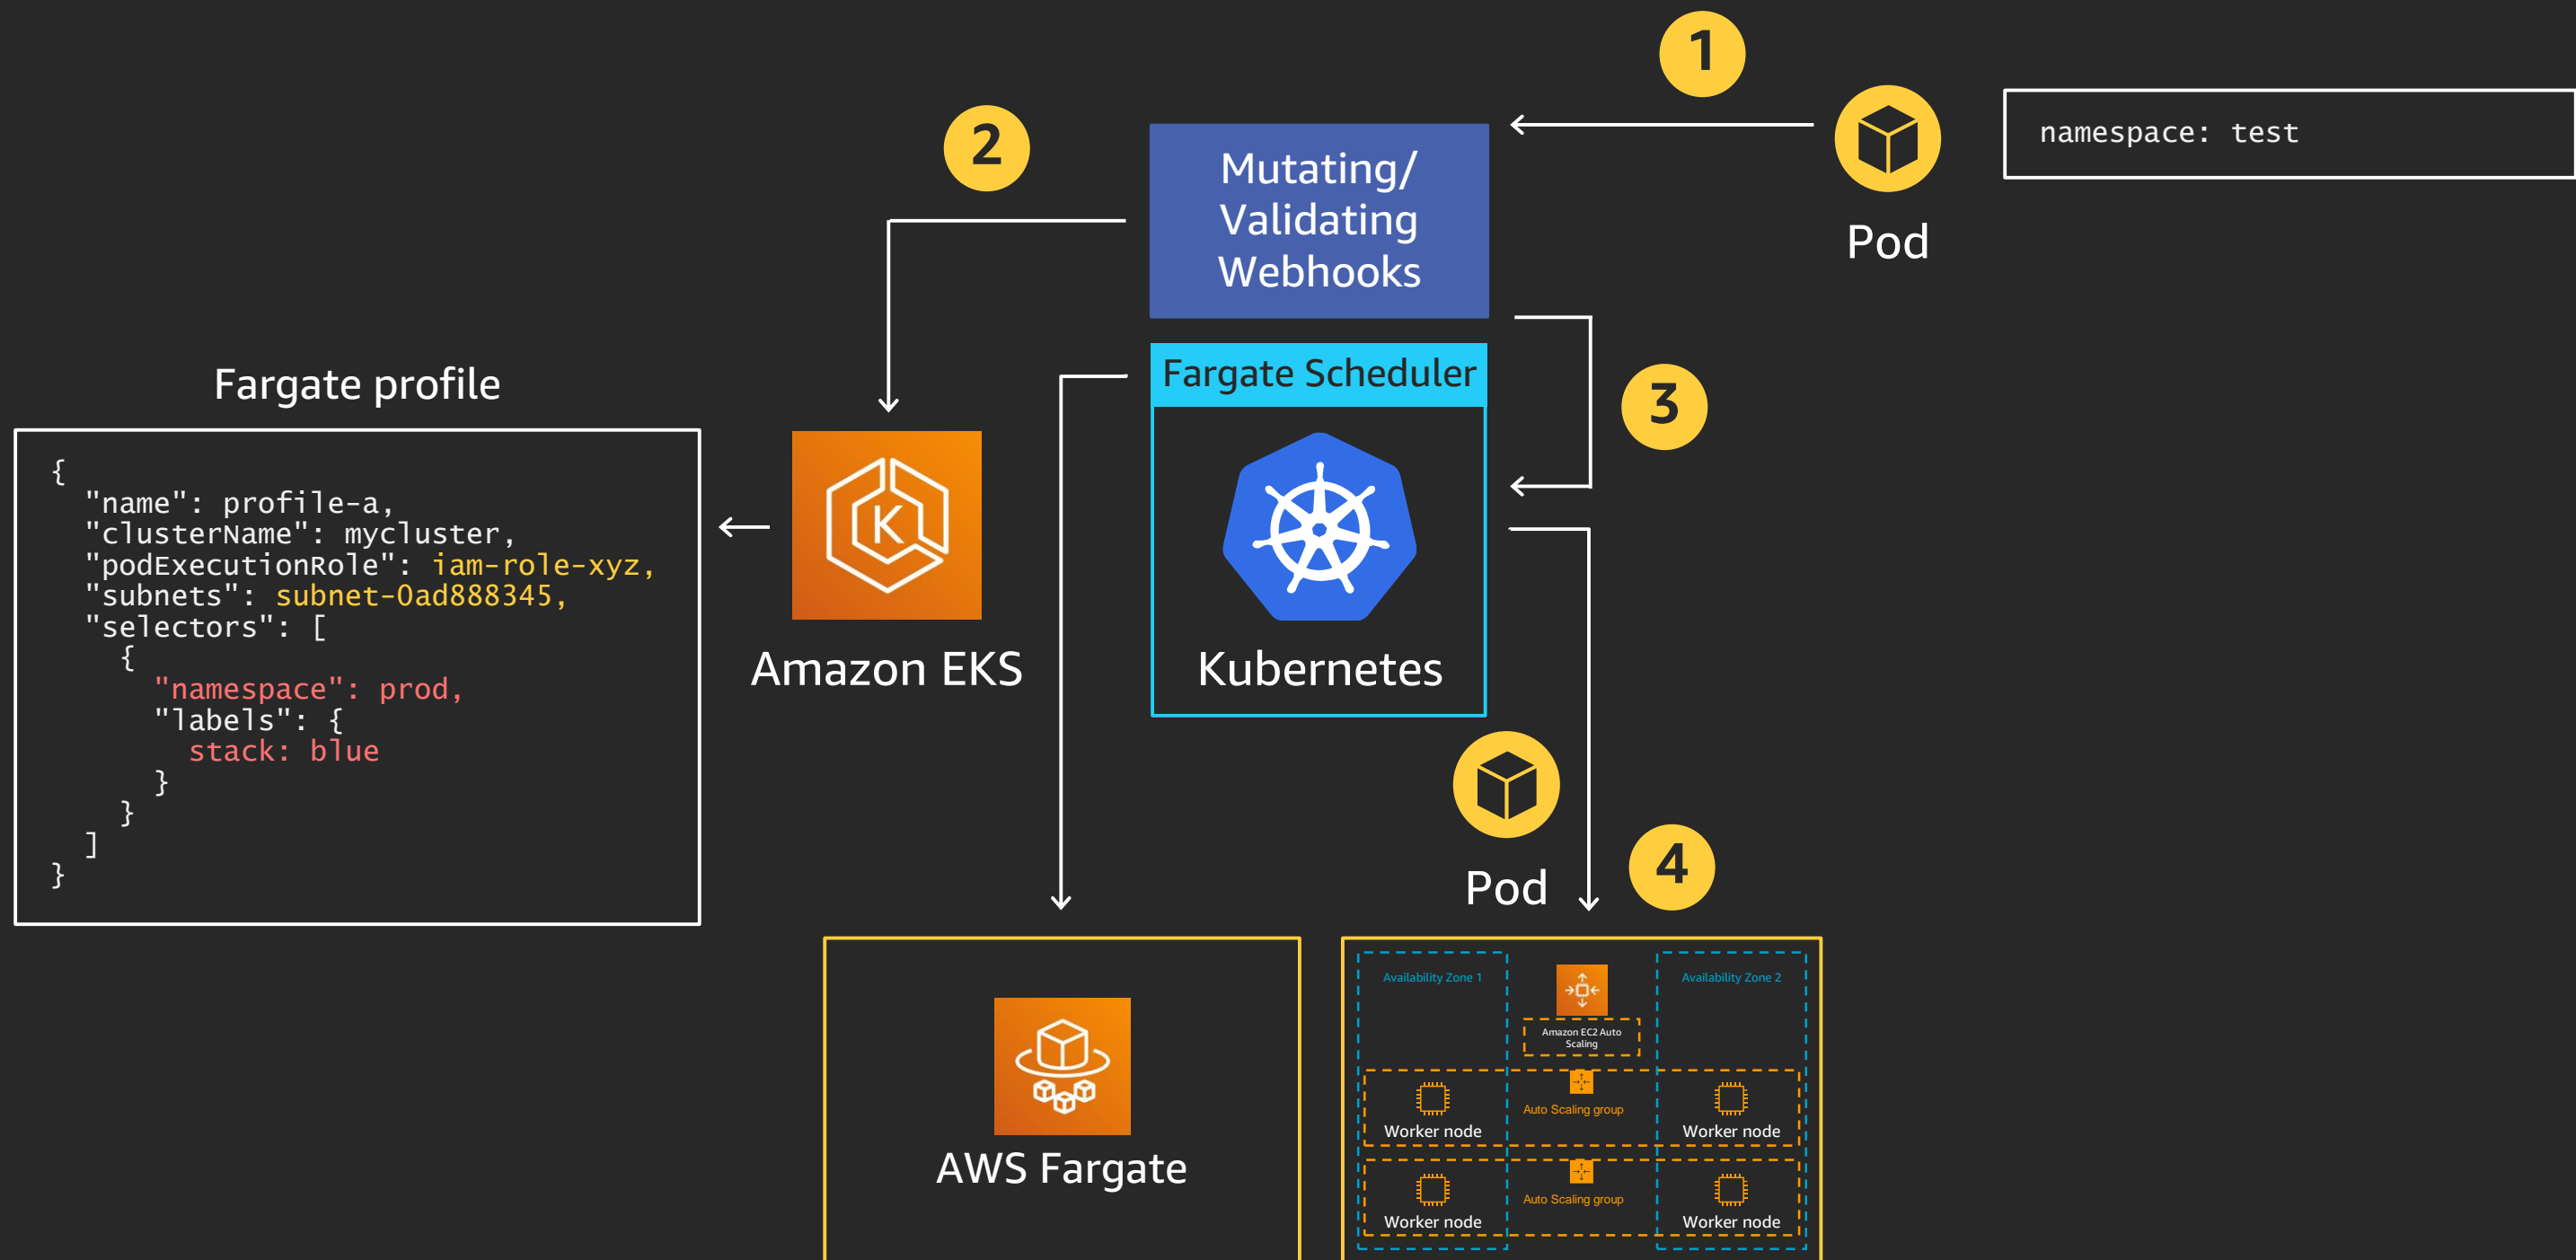
Parameters to "catch" the pod deployment

IAM Role to be associated to the kubelet

Simplified deployment flow



Simplified deployment flow




Need a custom pod spec to deploy to Fargate?

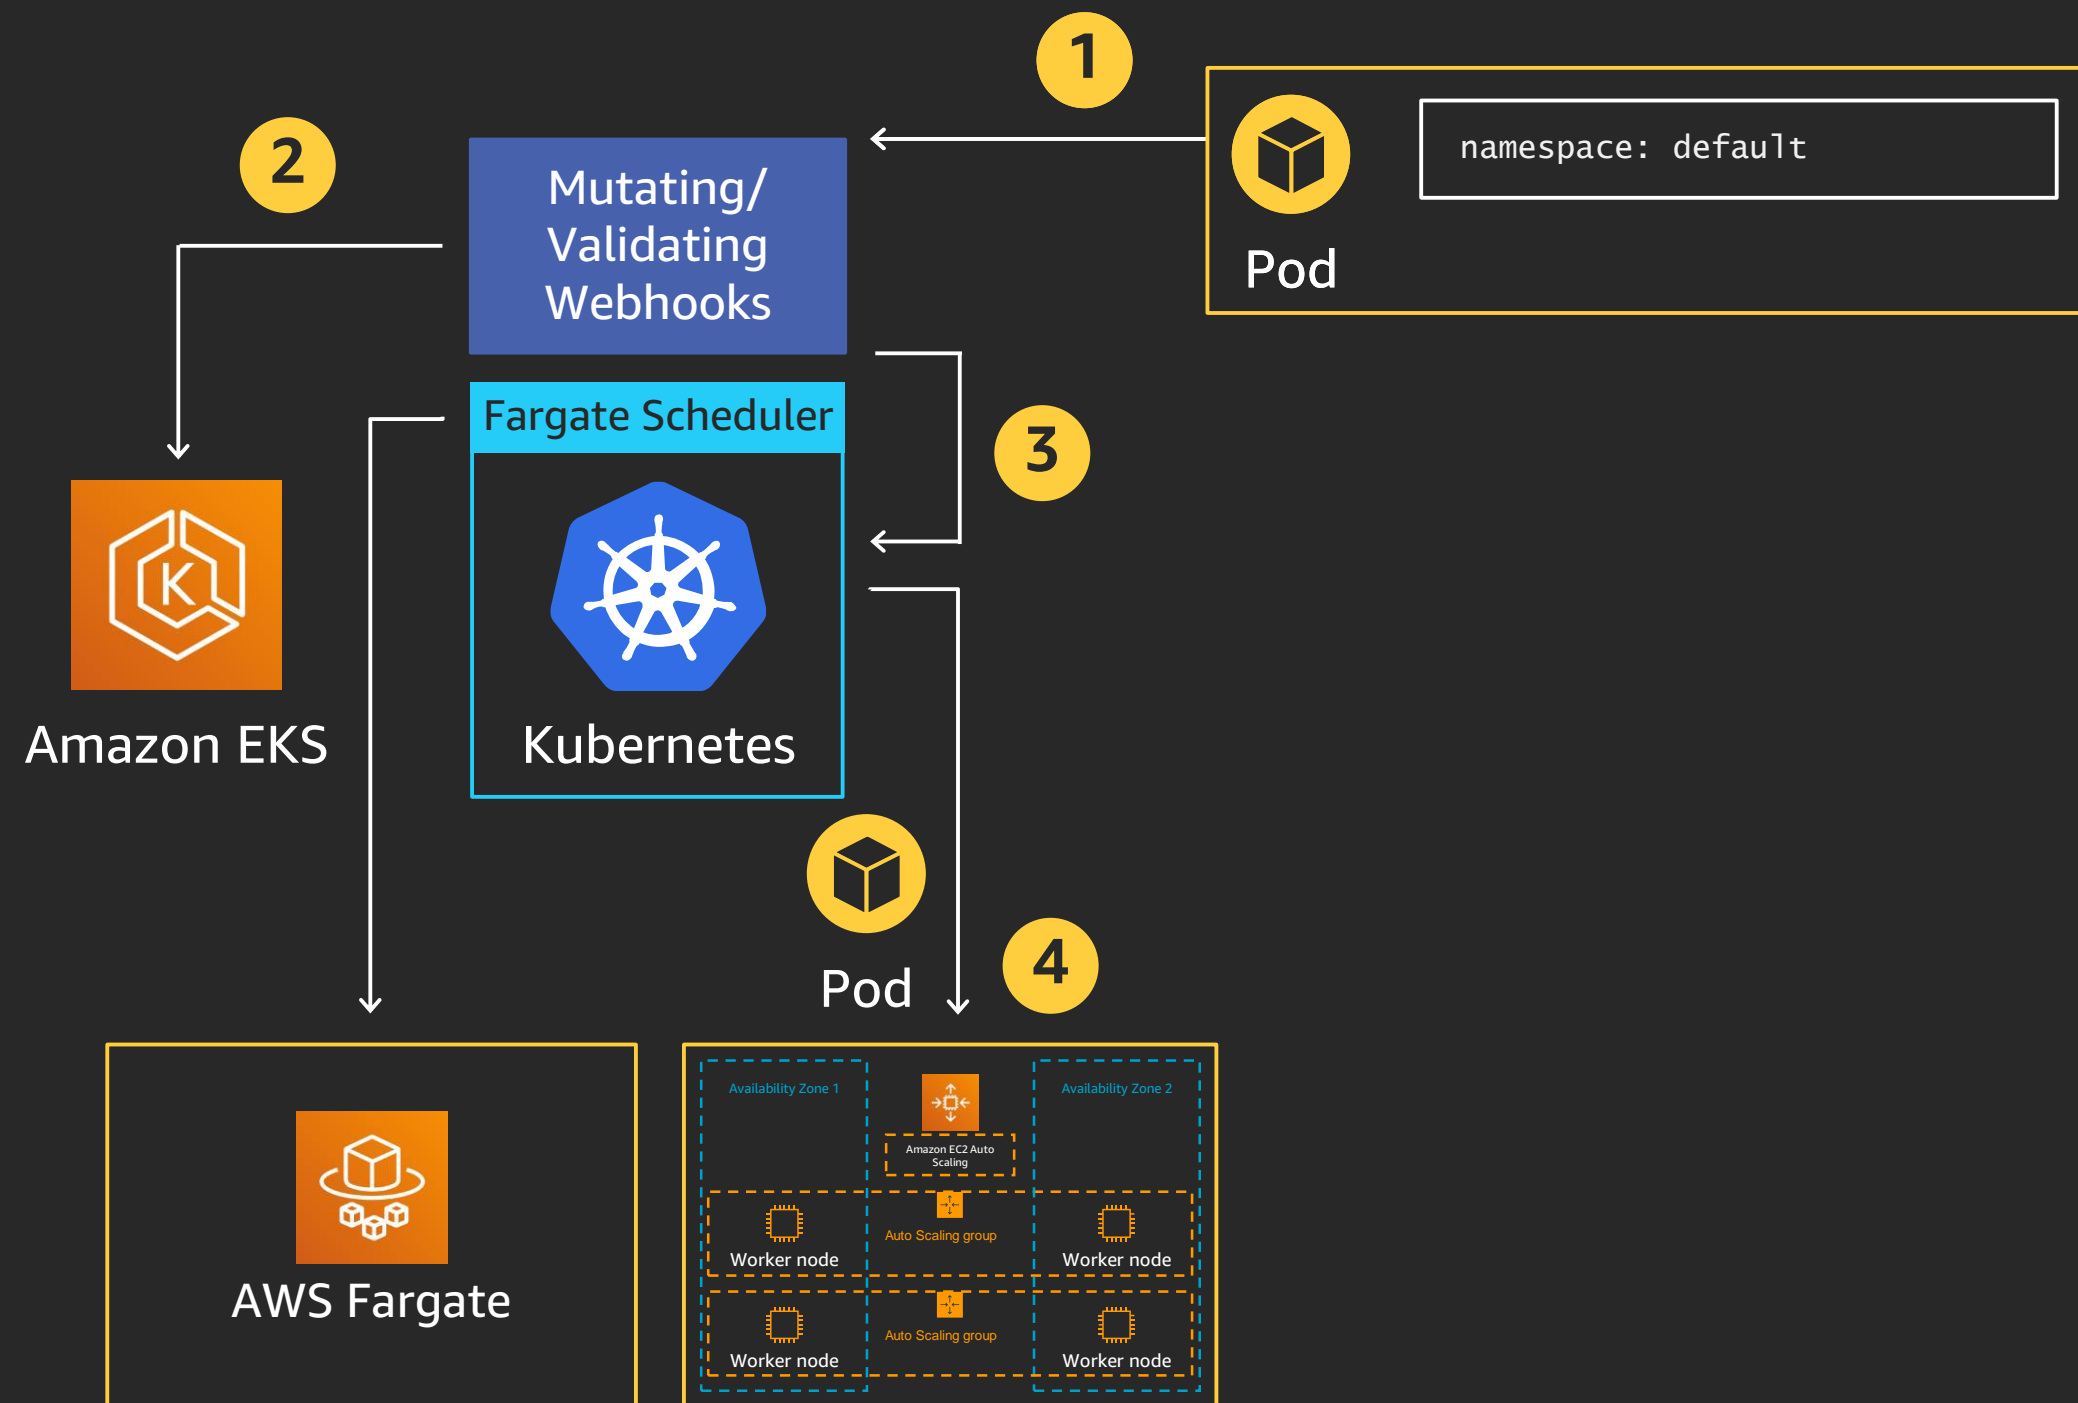
No

You can configure EKS to deploy to Fargate...
without touching your pod spec

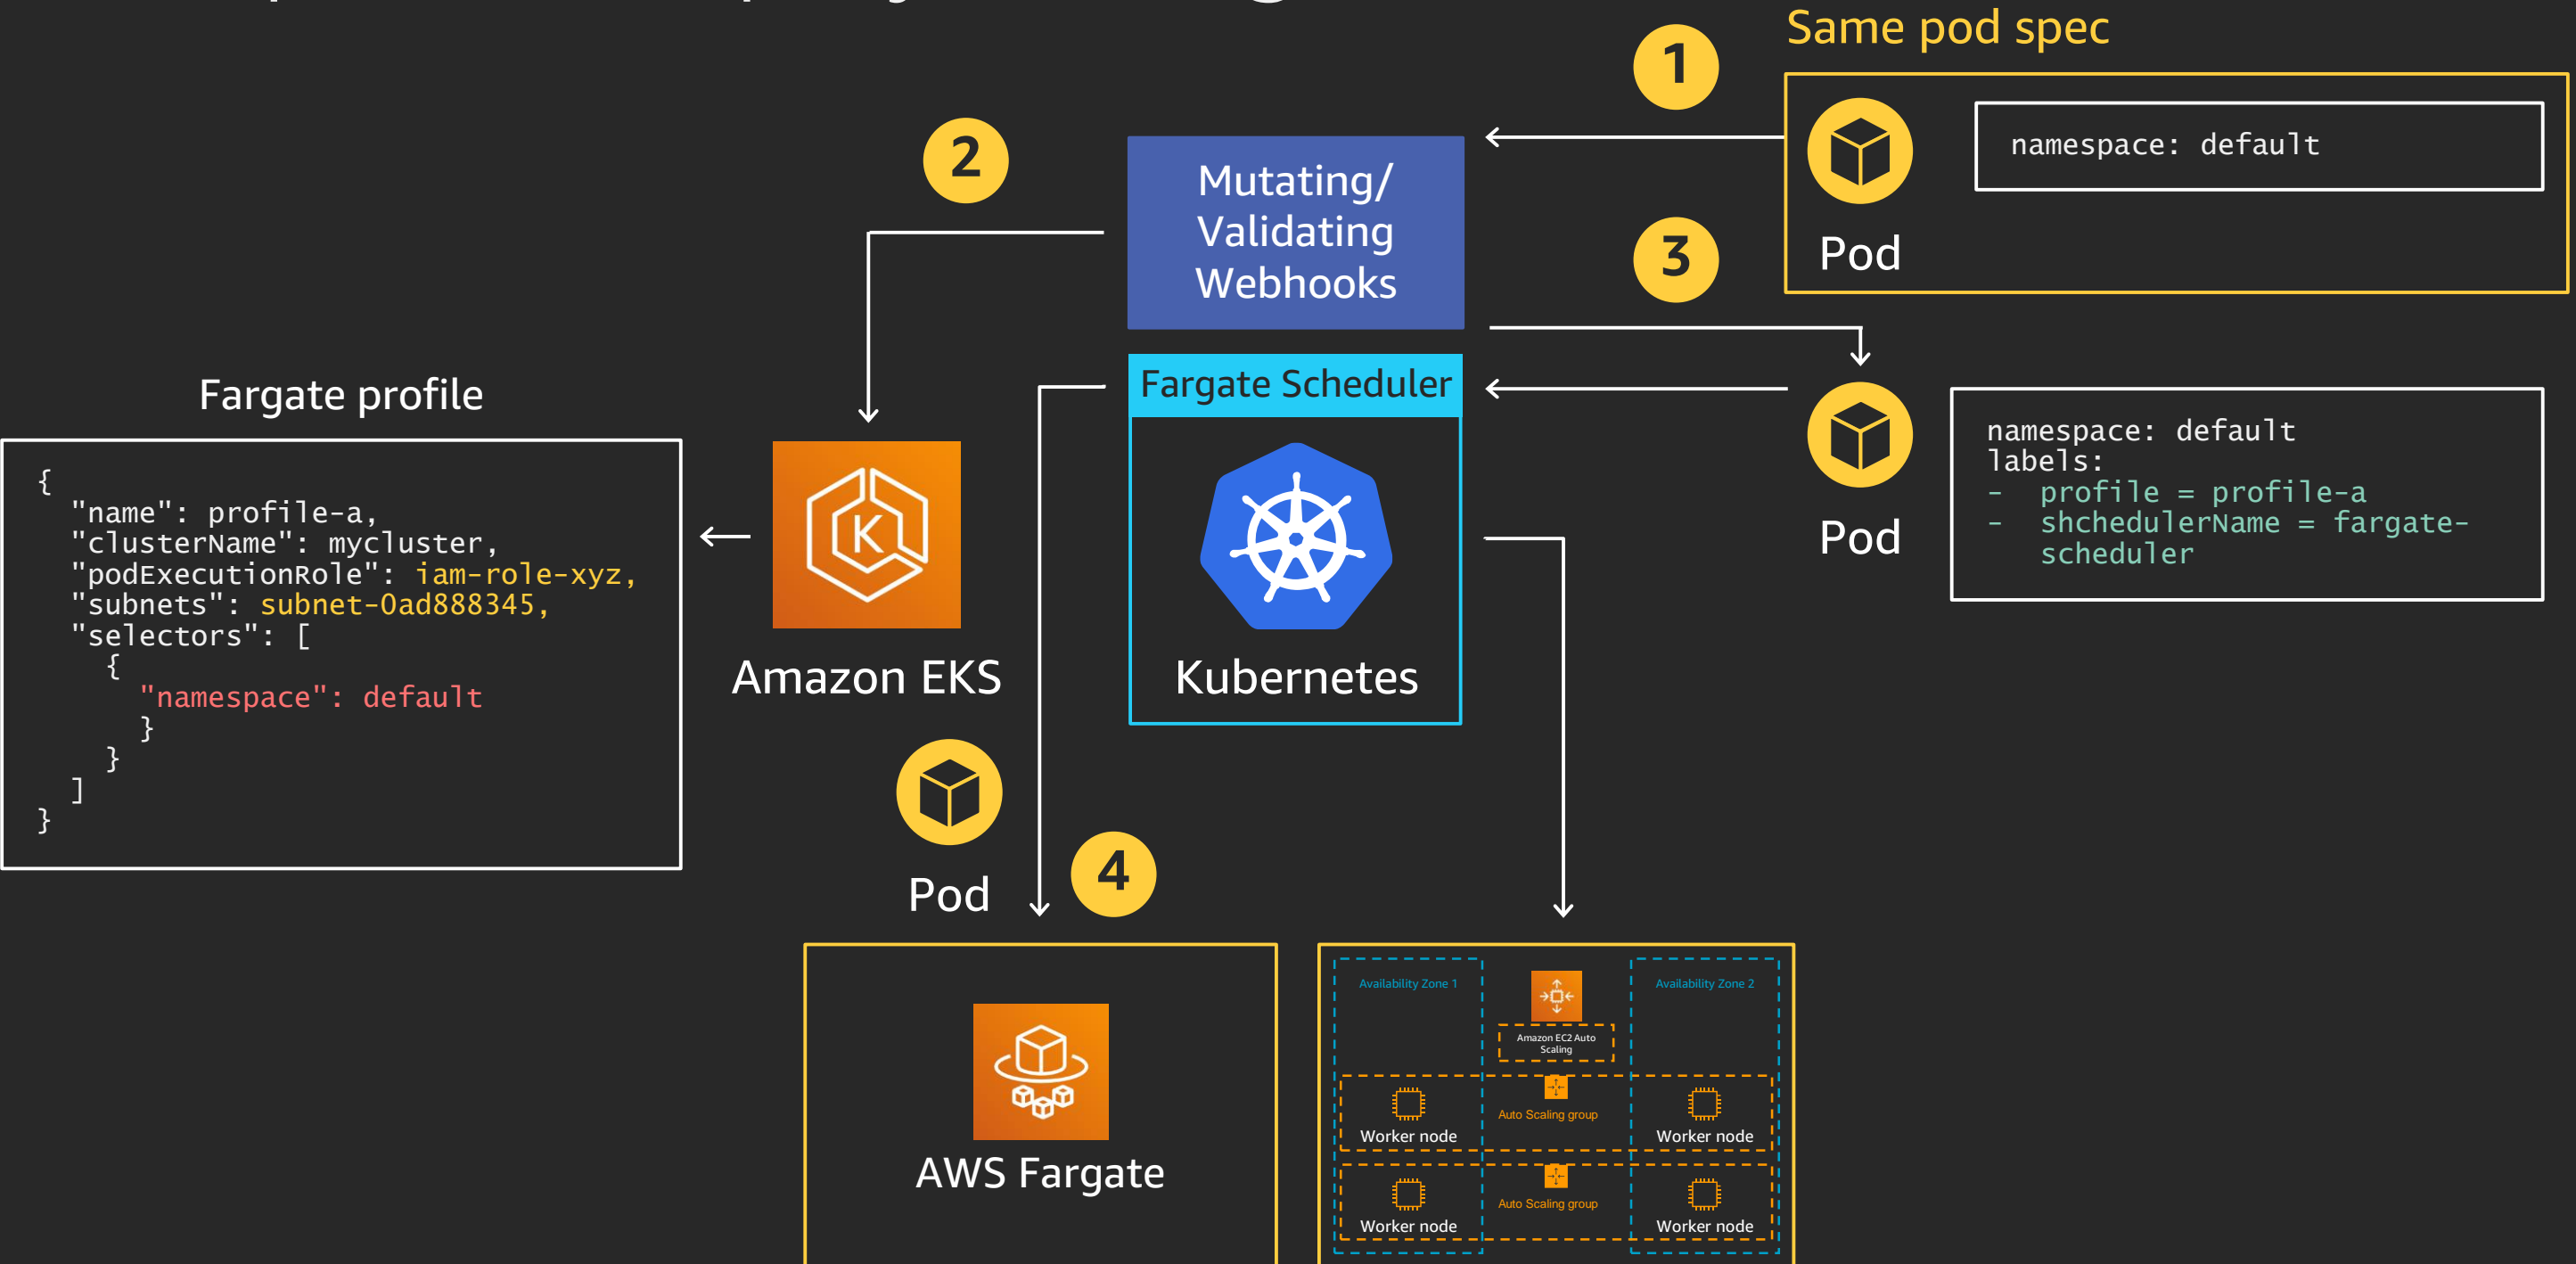
Example




Example: Deploy to workers




Example: Re-deploy to Fargate




How do we pick the size of the pod?




How do we pick the size of the pod?



How do we pick the size of the pod?




How do we pick the size of the pod?



How do we pick the size of the pod?

Another
example?



Init containers


Start sequentially and then stop

Containers

Long running

How do we pick the size of the pod?

Let's
consider the
requests only



Init containers

Start sequentially and then stop

Containers

Long running


How do we pick the size of the pod?

This

Let's add all
long-running
containers and
pick the biggest
number

Init containers
Start sequentially and then stop


Containers
Long running



How do we pick the size of the pod?

This

How do we
go from this
Pod config
example to a
Fargate size?



Init containers


Start sequentially and then stop

Containers

Long running

How do we pick the size of the pod?

This



CPU	Memory
256 (.25 vCPU)	512MB, 1GB, 2GB
512 (.5 vCPU)	1GB, 2GB, 3GB, 4GB
1024 (1 vCPU)	2GB, 3GB, 4GB, 5GB, 6GB, 7GB, 8GB
2048 (2 vCPU)	Between 4GB and 16GB in 1GB increments
4096 (4 vCPU)	Between 8GB and 30GB in 1GB increments

Fargate task size combinations

How do we pick the size of the pod?

This



Kubernetes components

CPU

Closest config
(rounded up)
is picked

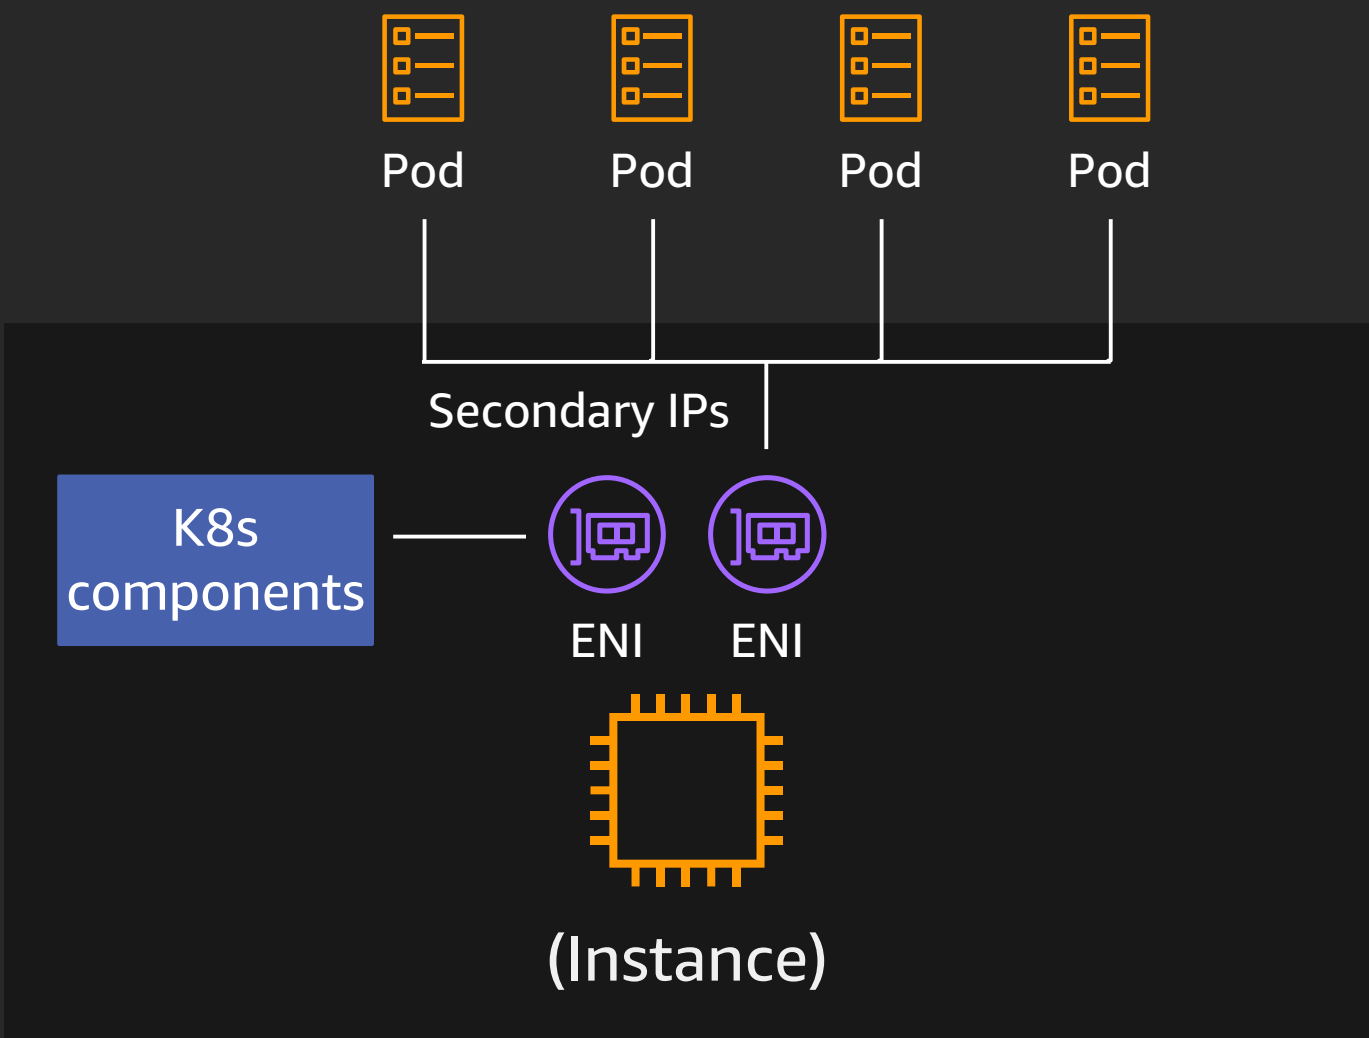


MEM

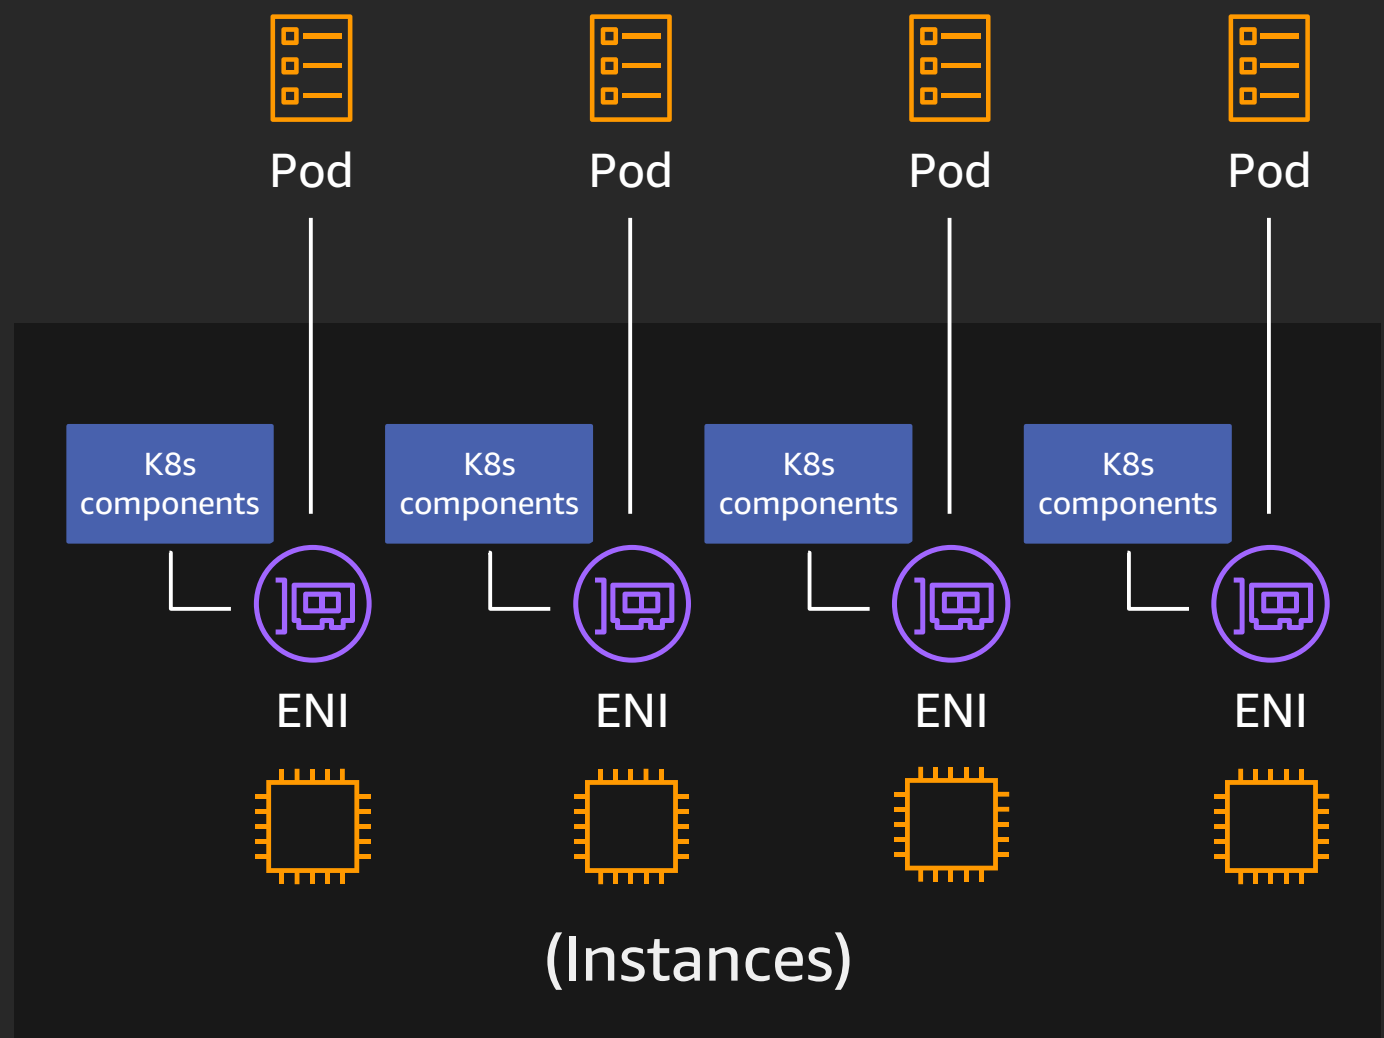
CPU	Memory
256 (.25 vCPU)	512MB, 1GB, 2GB
512 (.5 vCPU)	1GB, 2GB, 3GB, 4GB
1024 (1 vCPU)	2GB, 3GB, 4GB, 5GB, 6GB, 7GB, 8GB
2048 (2 vCPU)	Between 4GB and 16GB in 1GB increments
4096 (4 vCPU)	Between 8GB and 30GB in 1GB increments

Fargate task size combinations

Networking architecture




Worker nodes data plane




Fargate data plane

Security group considerations (at GA)

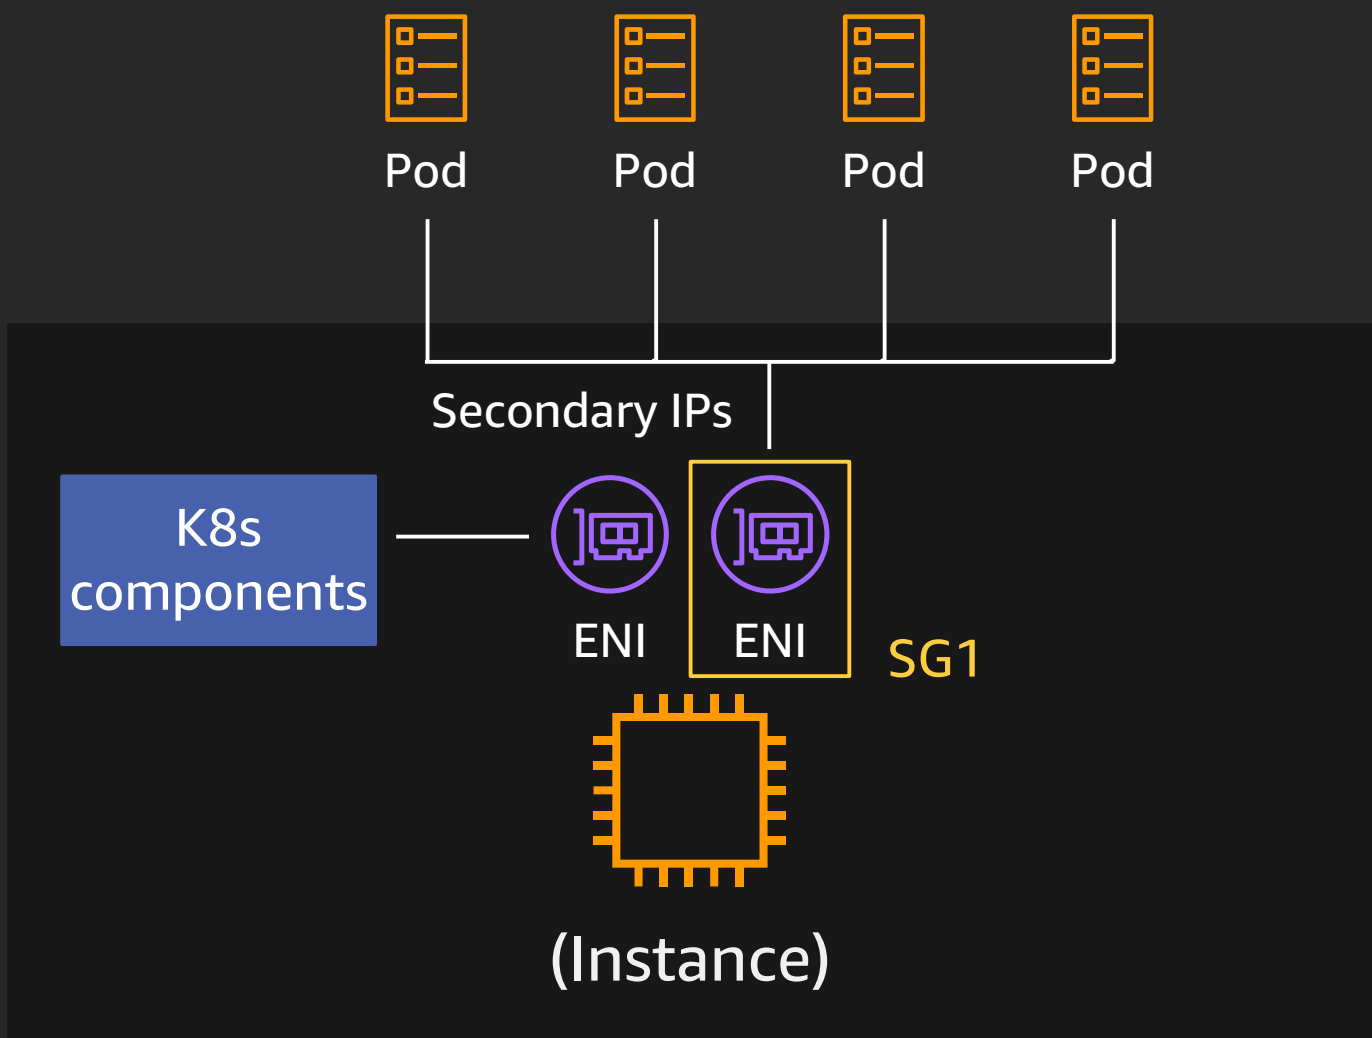


Worker nodes data plane

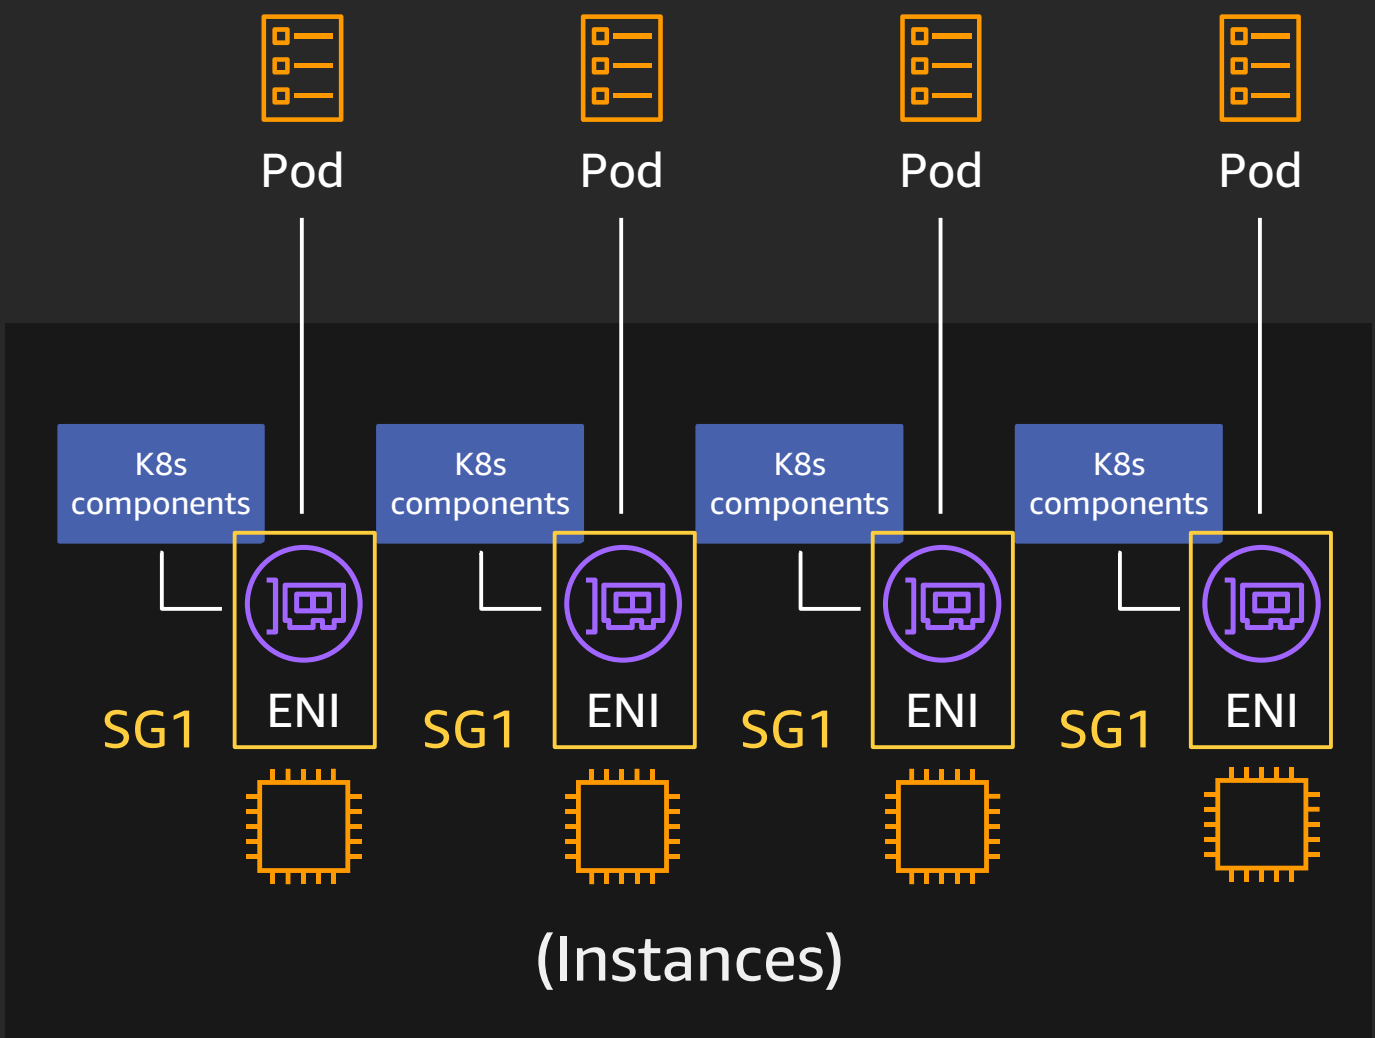


Fargate data plane

Security group considerations (at GA)



Worker nodes data plane



Fargate data plane

At GA the same SG will apply to ALL ENIs

Load Balancers considerations

ALB Ingress works as it normally does

Latest version includes the required code changes to make it work

NLB support is coming soon

Stay tuned

CLB will not work because it must target EC2 instances

There are no EC2 instances with EKS/Fargate

Storage options with EKS for Fargate

Fargate provides a local storage space for containers to share

This space is ephemeral and only lives for the time the pod lives

Persistent storage for Fargate is a frequent ask from customers

We are investigating the possibilities

EKSCTL support

EKSCTL supports Fargate and EC2 worker nodes

Make sure you use the latest version

It is possible to create a Fargate-only cluster

Or a combination of Fargate and managed node groups

EKSCTL takes care of some undifferentiated heavy lifting

Such as creating the Fargate profiles and more

```
> eksctl create cluster --fargate
```

Getting started

Recap: EKS for Fargate introduces UX changes

Things you no longer need to do

- ✔ Manage Kubernetes worker nodes
- ✔ Pay for unused capacity
- ✔ Use K8s Cluster Autoscaler (CA)

Things you get out of the box

- ✔ VM isolation at pod level
- ✔ Pod level billing
- ✔ Easy chargeback in multi-tenant scenarios

Things you can't do (for now)

- ✘ Deploy Daemonsets
- ✘ Use service type LoadBalancer (CLB/NLB)
- ✘ Running privileged containers
- ✘ Run stateful workloads

Limits: Things to keep in mind



AWS accounts have a soft limit of **100**
Fargate tasks/pods per region


You increase this limit



Due to the nature of the solution,
there's a **limit of 5,000 pods per cluster**

**K8s tests up to
5,000 workers per cluster**

Scalability: Things to keep in mind



Single individual pod start time may be longer on Fargate than on EC2

Each pod deployment sources a virtual node first from the Fargate fleet

Pod deployments at scale may be faster due to Fargate parallelism

E.g., think of the delay that Cluster Autoscaler can introduce in sourcing new EC2 capacity

Pricing

Standard EKS cluster pricing **\$0.20 per hour**

Standard Fargate Pricing for vCPU and memory

Availability

Available today for all new 1.14 clusters

- Create a new cluster
- Update a 1.13 cluster to 1.14
- We'll automatically update existing 1.14 clusters in the coming weeks

Use EKS with Fargate in

- Virginia (us-east-1)
- Ohio (us-east-2)
- Dublin (eu-west-1)
- Tokyo (ap-northeast-1)

Thank you!

Nathan Taber

@nctaber

Massimo Re Ferrè

@mreferre



Please complete the session
survey in the mobile app.